



Amirkabir University of Technology
(Tehran Polytechnic)
Department of Computer Engineering and Information Technology

Processes (فرآیندها)

Hamid R. Zarandi

h_zarandi@aut.ac.ir

Definition

مختصرات OS، بـ اجری برخواهی کار *

➤ Process

- A program in execution; process execution must progress in sequential fashion
 - In time-sharing sys: unit of work → *عملیاتی کار*

- All processes are executed **concurrently**

➤ Process vs. Job?

- **Passive:** program
- **Active:** process
 - Program becomes **process** when executable file loaded into **memory**
 - One program can be several processes

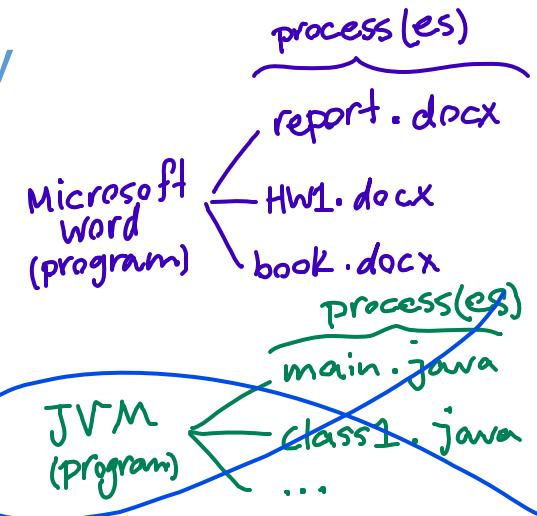
- **Question?**

- java program

Parallel → *عملیاتی*
 concurrently → *عملیاتی هم‌زمان*

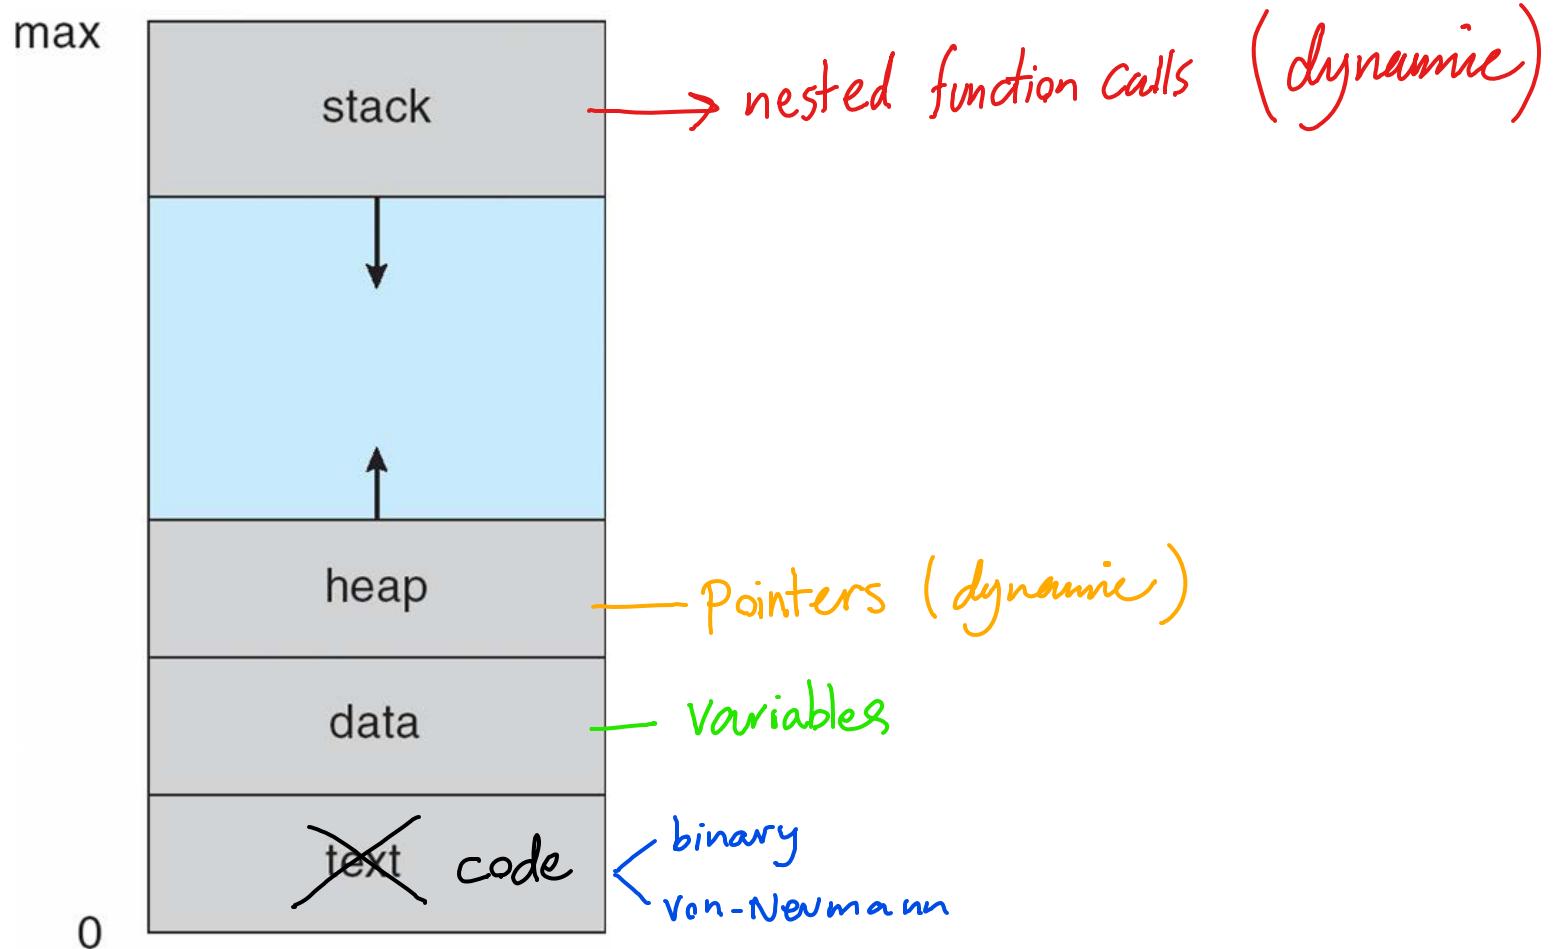
program, process *کار* : one-to-many

Runnable ← *نیازمند اجرا*, *کار*



Process in memory

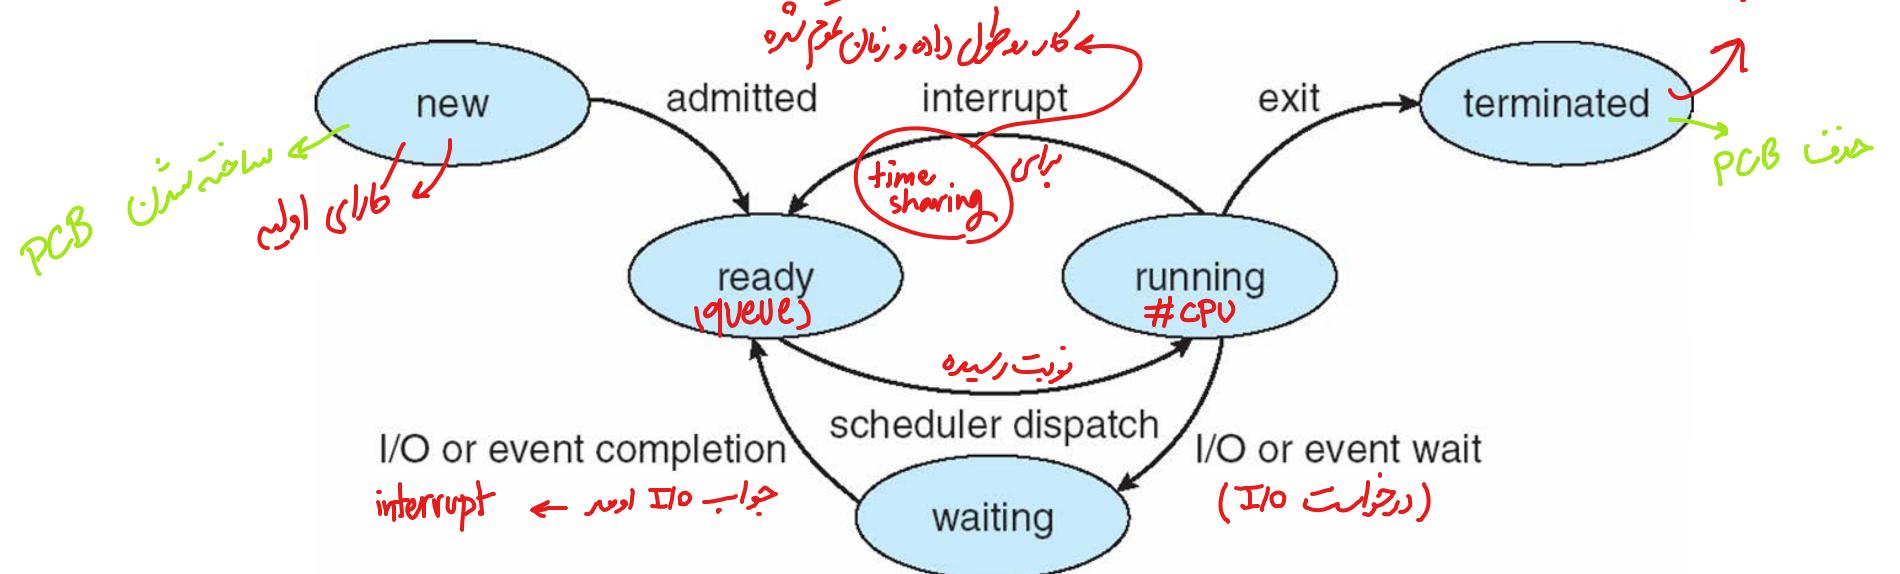
program \rightarrow Process (ج)
 ، heap, data \rightarrow {text \rightarrow
 . . . \rightarrow stack}



Process state

➤ As a process executes, it changes **state**

- **new**: The process is being created → *ایجاد کریج اولین / هنوز PC درون اجرا نمی‌شود*
- **running**: Instructions are being executed → *PL(دستور) در CPU / Van-Neumann*
- **waiting**: The process is waiting for some event to occur → *using I/O / خروجی مکانیکی*
- **ready**: The process is waiting to be assigned to a processor → *در CPU، اینجا صفت ایجاد شده است*
- **terminated**: The process has finished execution → *MM(منابع) + بین خالی از ربوط + جزء اخراجی / (Ready Queue)*

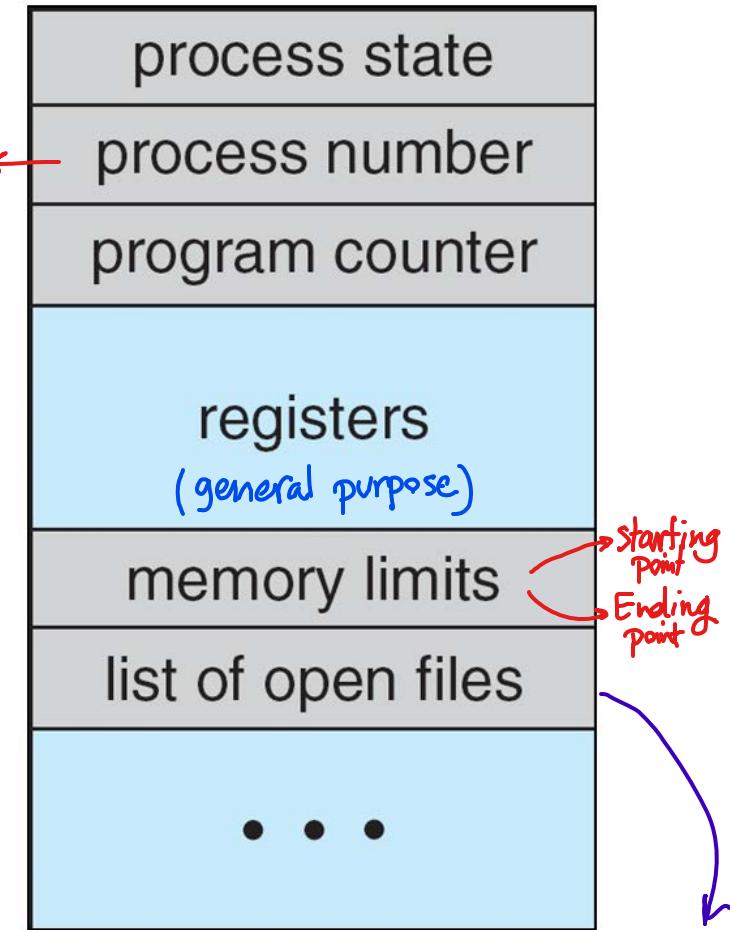


Process Control Block (PCB)

- How to manage processes?
- Information associated with each process
(also task control block) → Linux

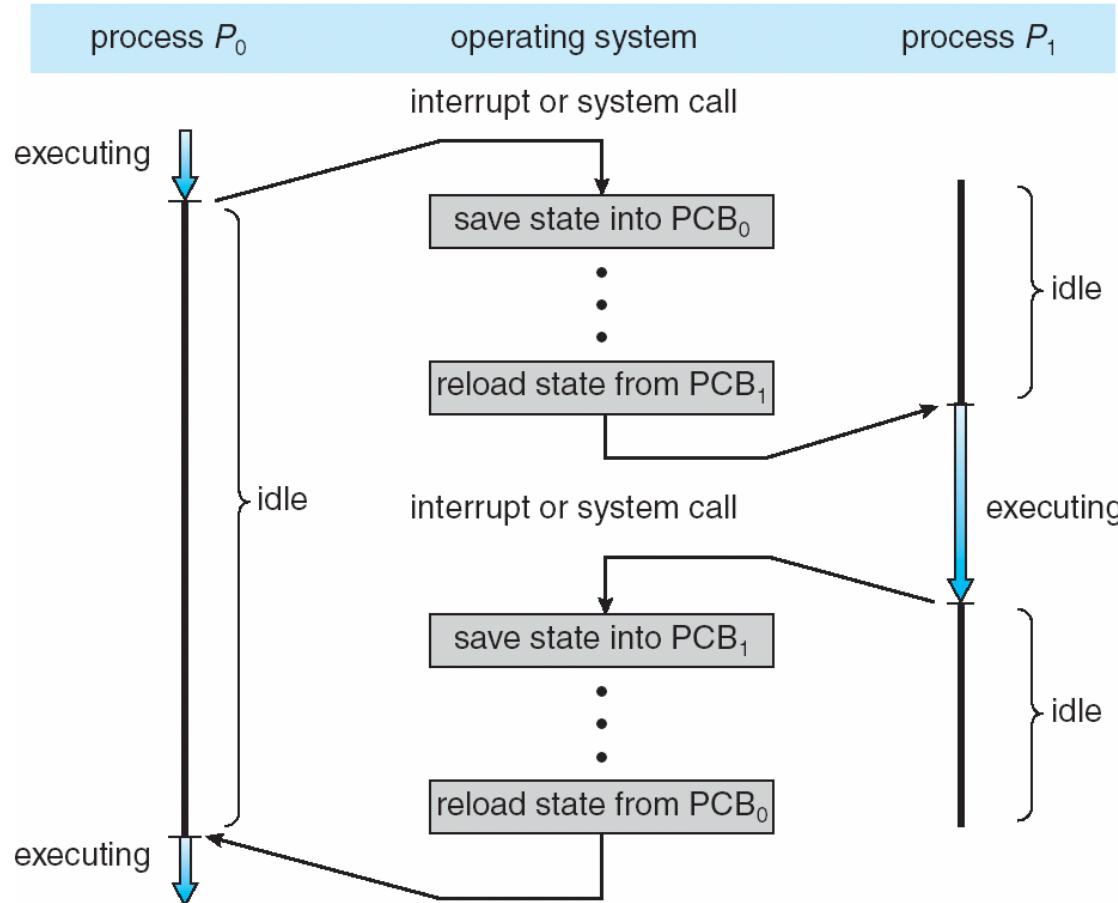
- Process state
- Program counter
- CPU registers – contents of all process-centric registers
- CPU scheduling info. – priorities, scheduling queue pointers
- Memory-management info. – memory allocated to the process
- Accounting info. – CPU used, clock time elapsed since start, time limits
- I/O status info. – I/O devices allocated to process, list of open files

unique ID



"This file is open in another program"

CPU switch from process to process



عامل رفع از OS ~ process
 interrupt : عامل خارجی
 system call : داخلي

- ID نه PCB هر *

زنگ بروط برخوبی link ایند

HLT از idle شدن Δ

میتوانید OS را اجرا میکنید.

switch کردن بین برنامه *

پیاز به OS میگیریم.

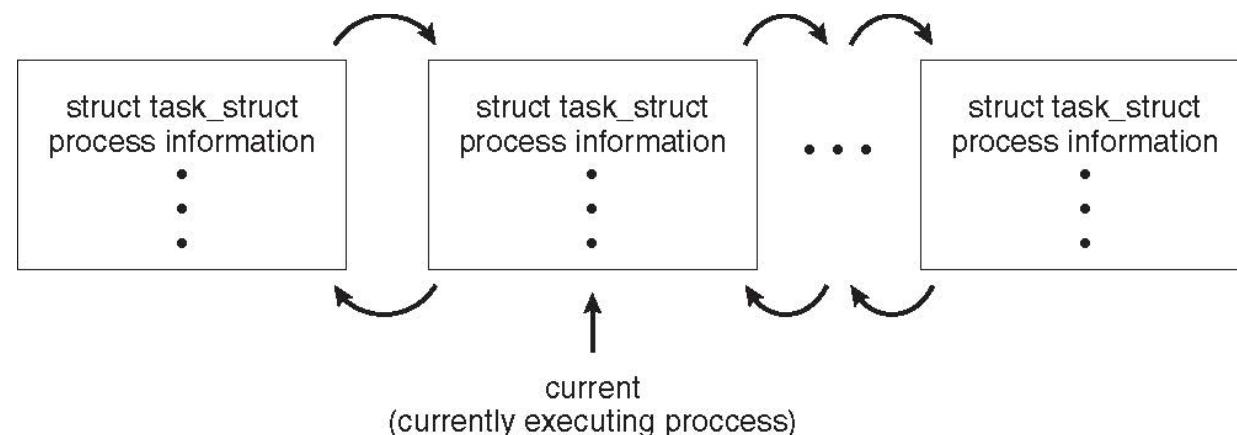
Process representation in Linux

Represented by the C structure task_struct

کوئن میل (Queue)
to process

$O(n)$ ←

```
pid t_pid; /* process identifier */  
long state; /* state of the process */  
unsigned int time_slice /* scheduling information */  
struct task_struct *parent; /* this process's parent */  
struct list_head children; /* this process's children */  
struct files_struct *files; /* list of open files */  
struct mm_struct *mm; /* address space of this process */
```



Process scheduling

- **Process scheduler selects among available processes for next execution on CPU**

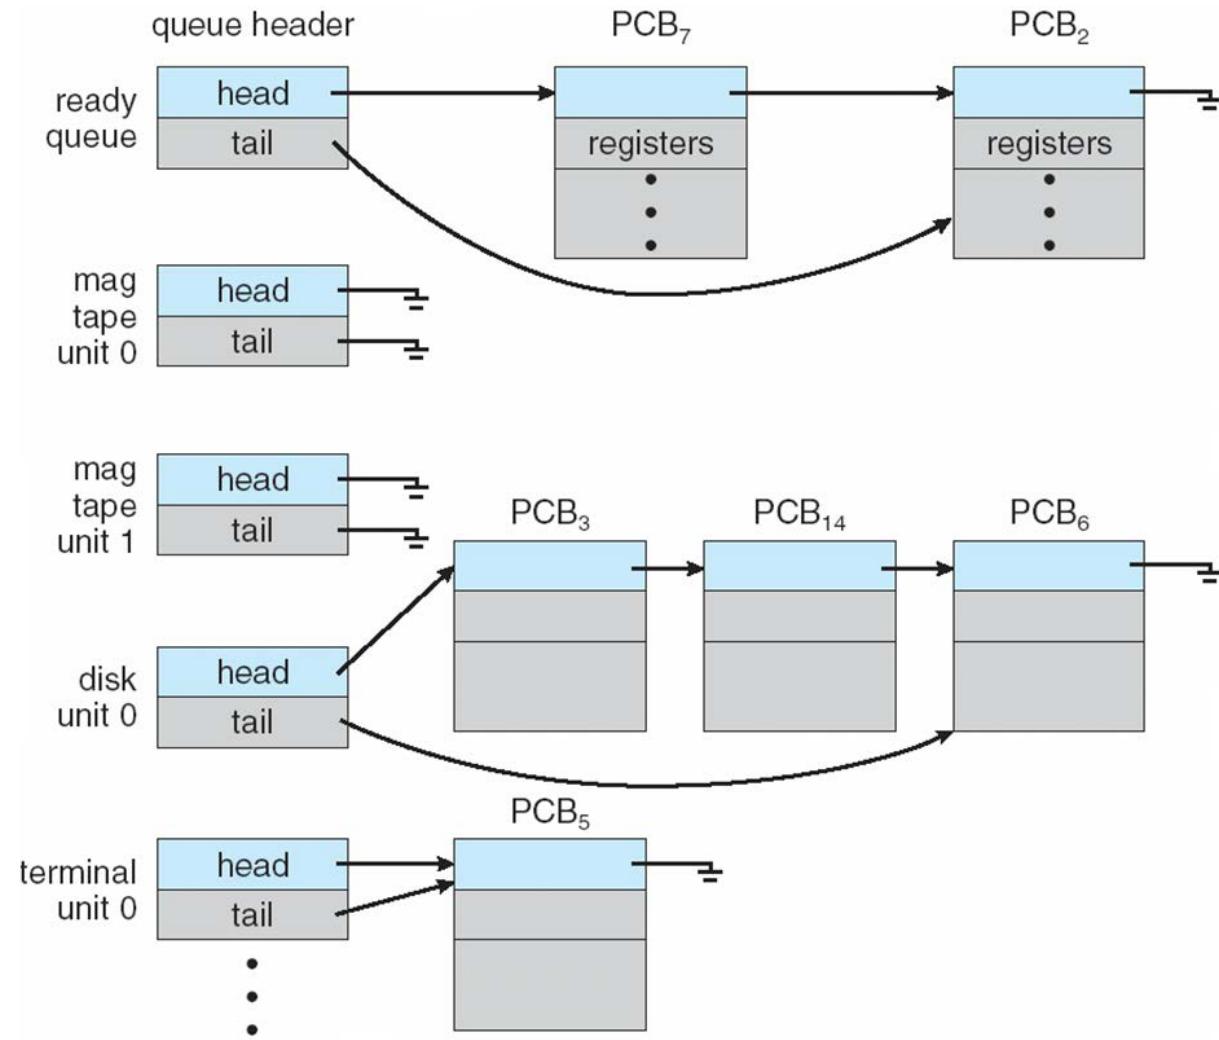
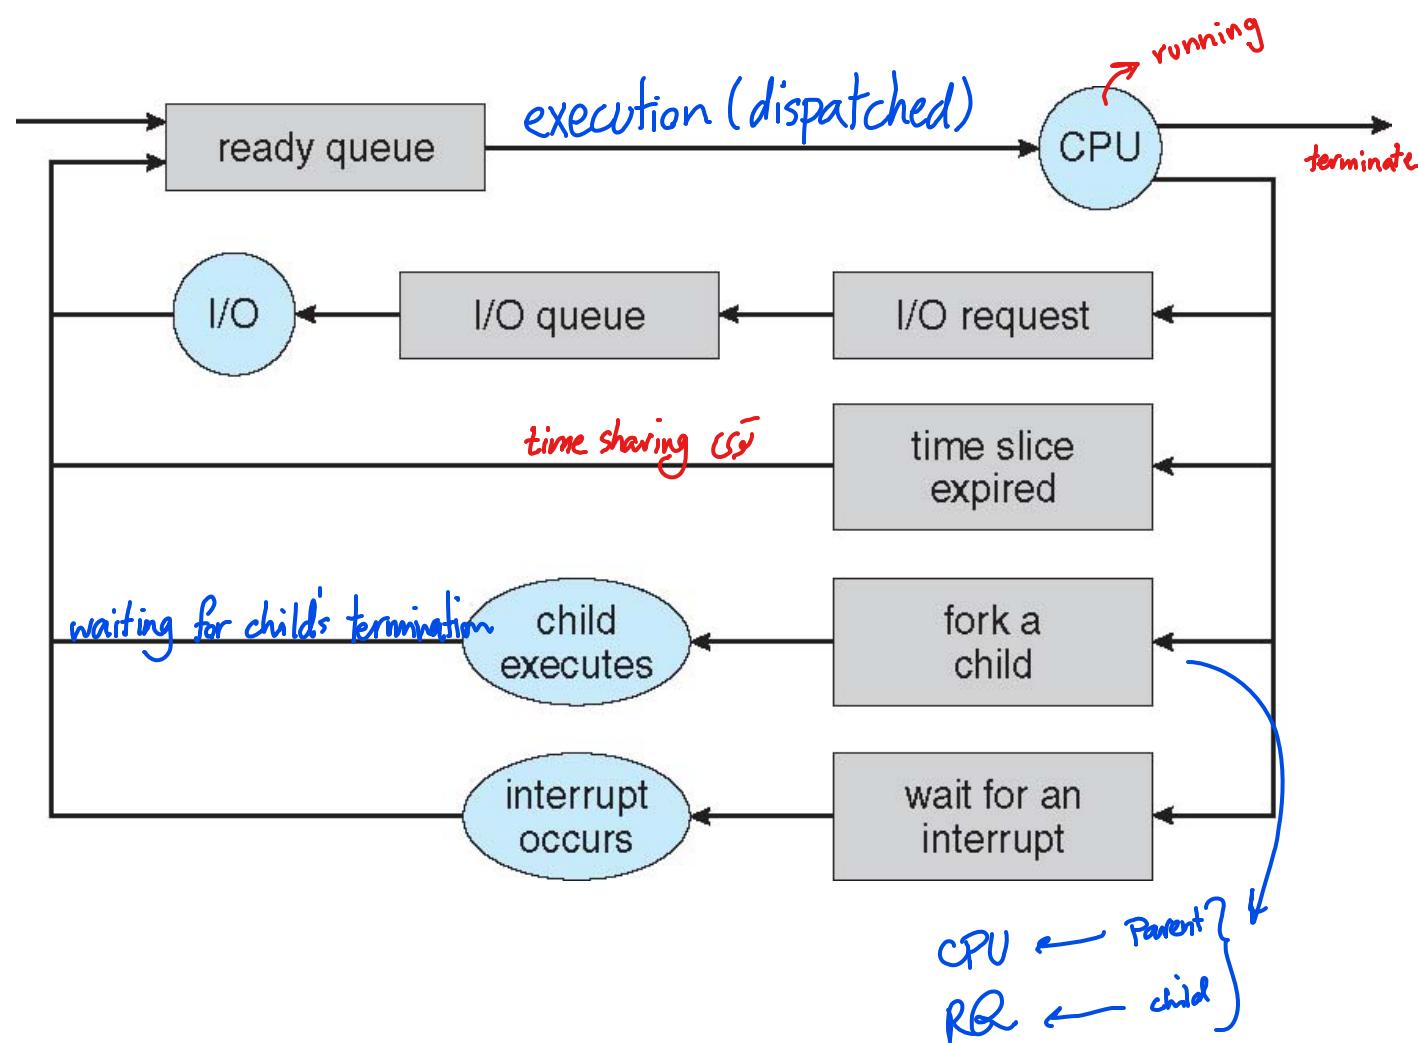


Diagram representation of process scheduling

- **Queueing diagram**
represents queues,
resources, flows

- {  : queues
- : resources that serve the queue



Schedulers



➤ Short-term scheduler (or CPU scheduler) سیک / وجہ الرانی / ready queue نتایج برقرار رکھنے والے CPU مائن بائیم

- selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (milliseconds) \Rightarrow (must be fast)

➤ Long-term scheduler (or job scheduler) کسی / وجود غیر رانی / ready queue جوں جتن کر دیں

- selects which processes should be brought into the ready queue
 - Long-term scheduler is invoked infrequently (seconds, minutes) \Rightarrow (may be slow)
 - The long-term scheduler controls the degree of multiprogramming

➤ Processes:

- I/O-bound
 - spends more time doing I/O than computations, many short CPU bursts
- CPU-bound
 - spends more time doing computations; few very long CPU bursts

➤ Long-term scheduler strives for good *process mix*

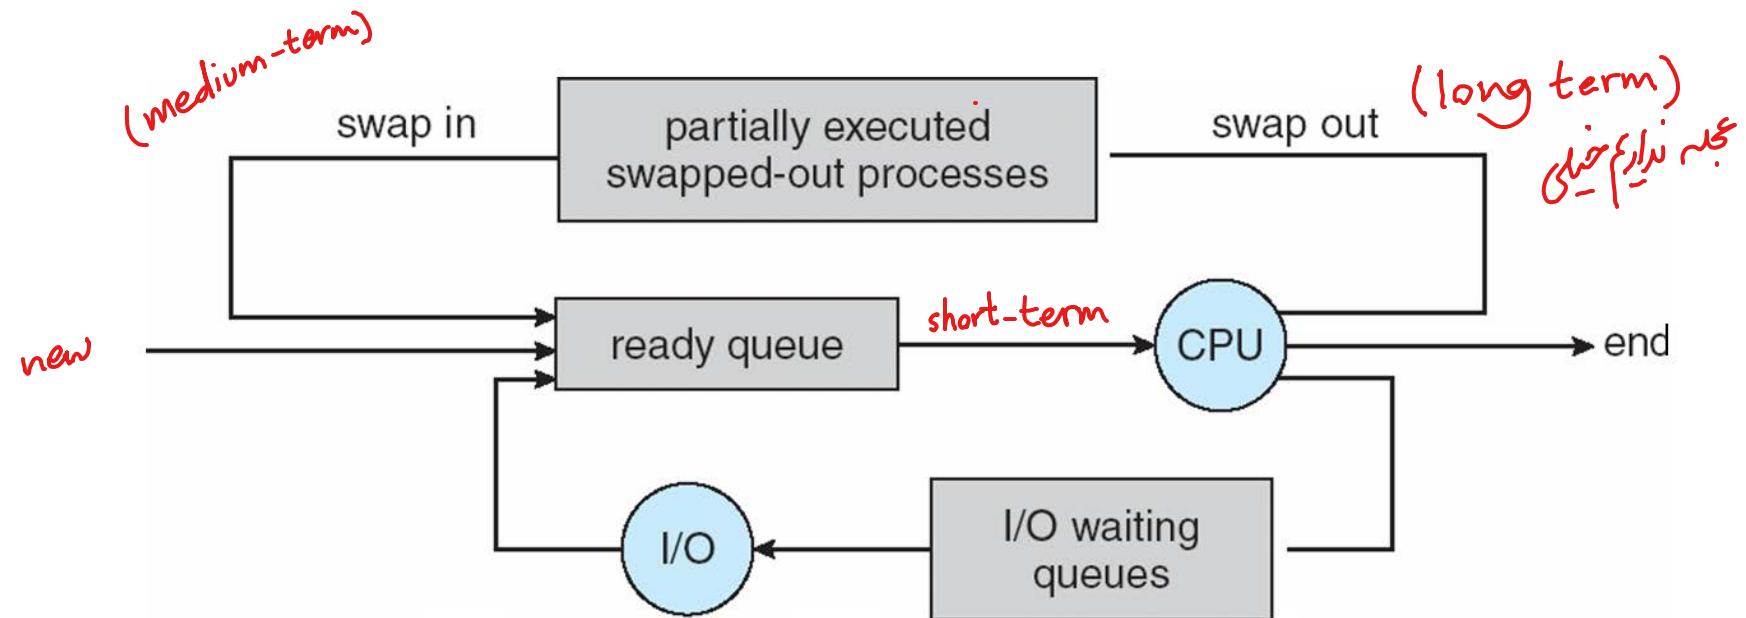
time sharing systems (UNIX - Microsoft windows) have no long-term scheduler but put every new process in Mem. for short-term.

Example of standard API

new process, Ready queue, long-term queue

➤ Medium-term scheduler

- Can be added if degree of multiple programming needs to decrease
- Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**



Context switch (تعویض متن)

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**

- Context** of a process represented in the PCB

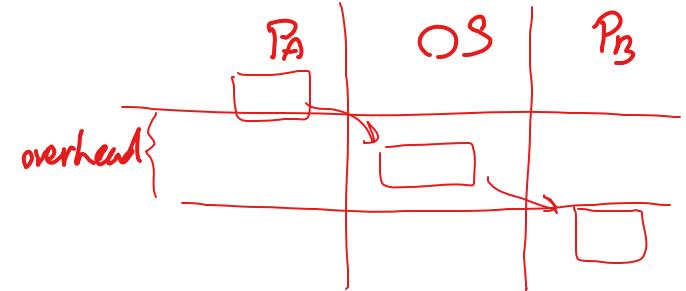
- Context-switch time is **overhead**; the system does no useful work while switching

 - The more complex the OS and the PCB → the longer the context switch

- Time dependent on hardware support

 - Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once

\uparrow memory speed
 # of regs must be copied
 instructions



\downarrow overhead (هزینه) \Leftarrow تا فرایند نئی

Process creation

➤ Parent vs. Child

(int)

➤ Generally, process identified and managed via a process identifier (pid)

→ variables, memory allocations, ...

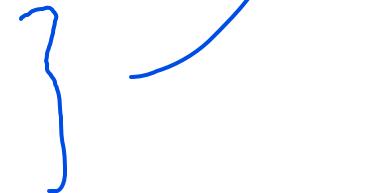
➤ Resource sharing options

- Parent and children share all resources
- Children share subset of parent's resources
- Parent and child share no resources

in (int) parent b/w

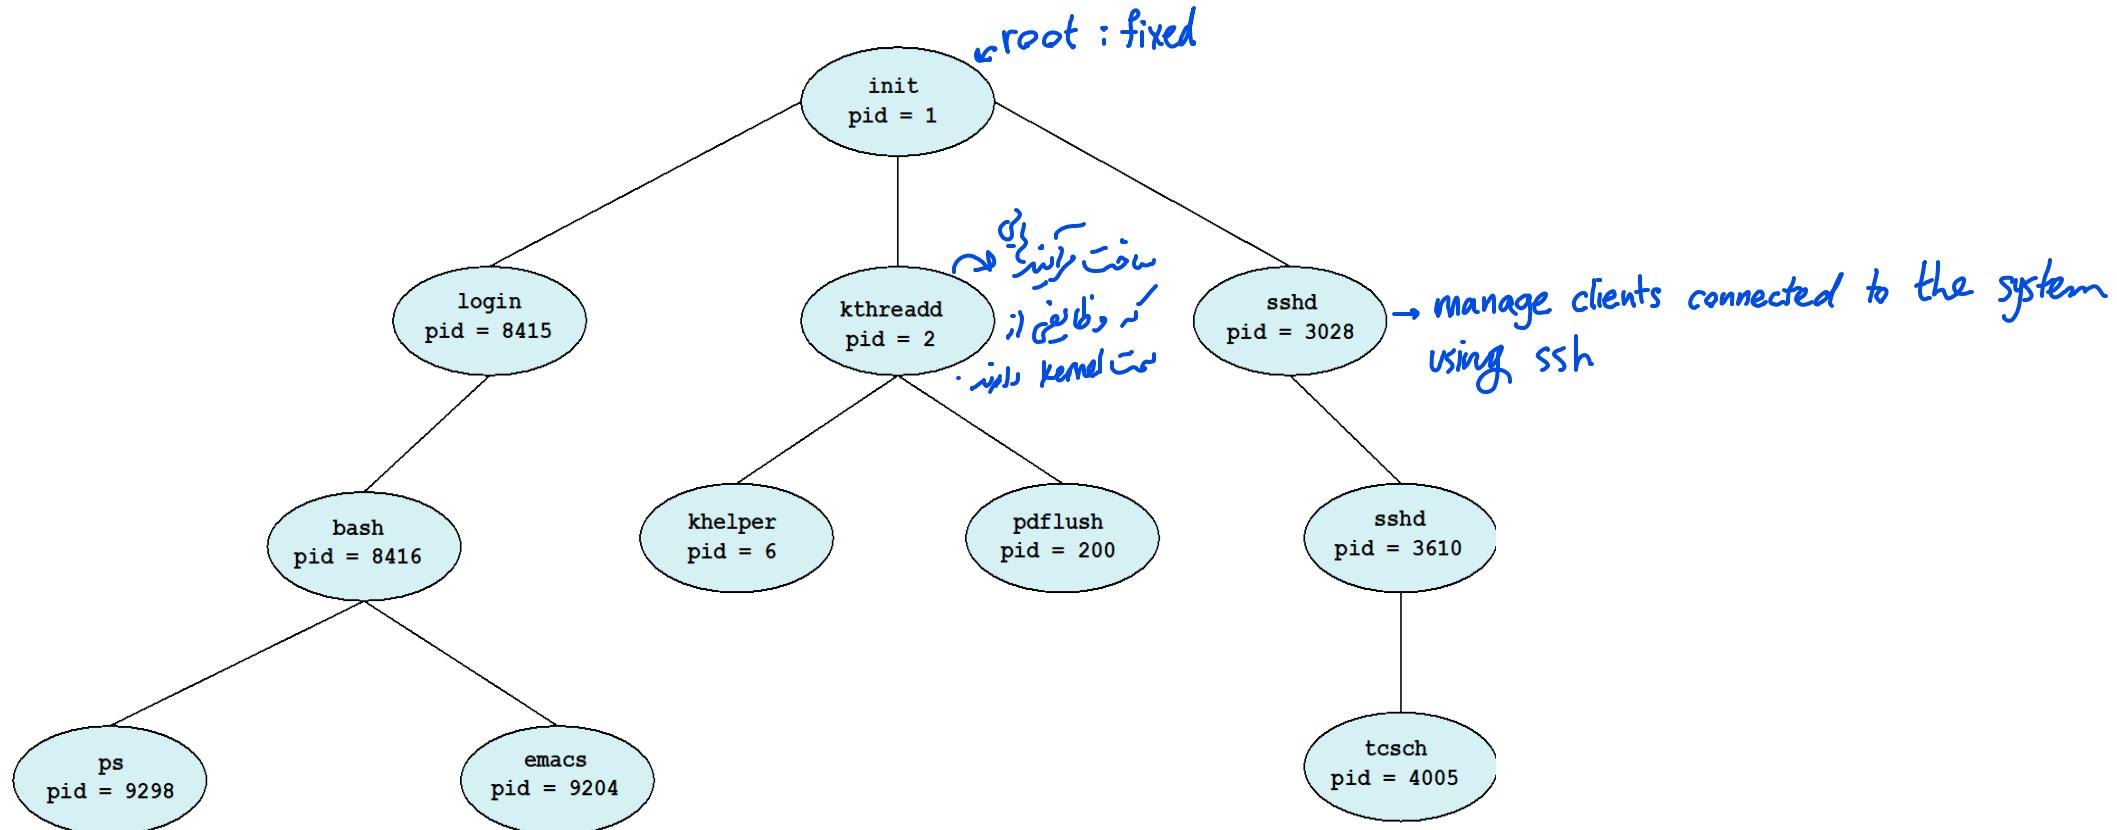
➤ Execution options

- Parent and children execute concurrently
- Parent waits until children terminate



A tree of processes in Linux

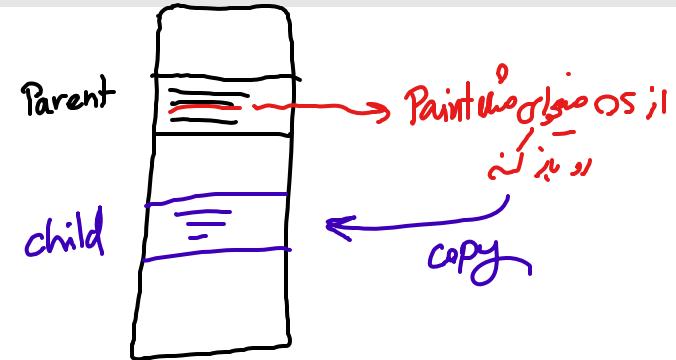
(لوداً بَوْصَارَهُ نَرَسَ) (كِبَيْ حَوْدَهُونَ) (جَنْجِي)



Process creation

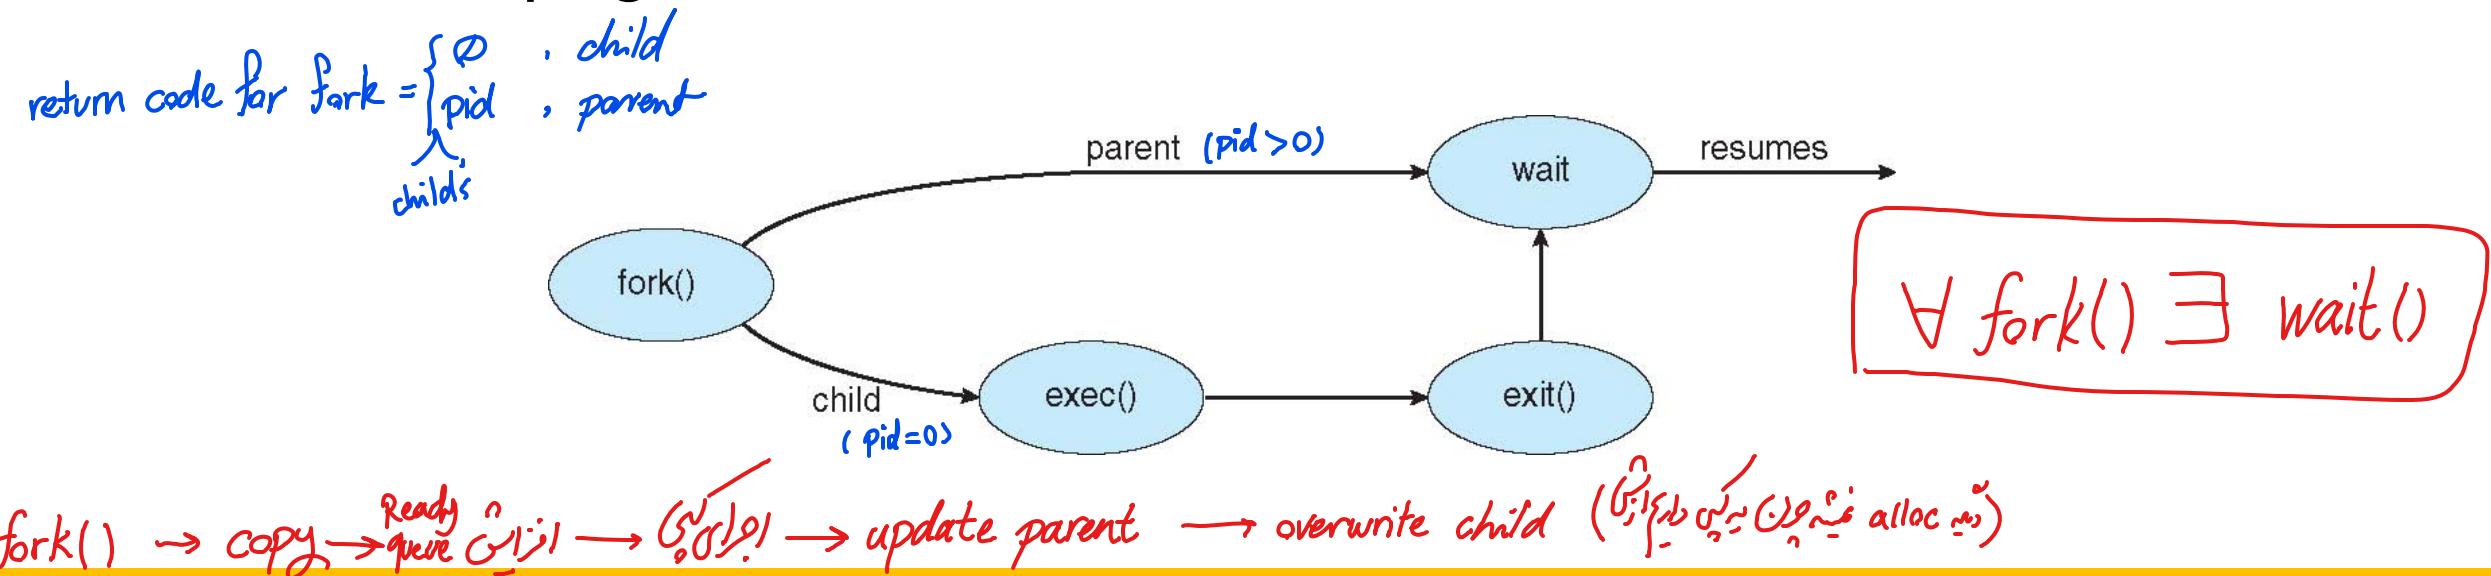
➤ Address space

- Child duplicate of parent
- Child has a program loaded into it



➤ UNIX examples

- **fork()** system call creates new process
- **exec()** system call used after a **fork()** to replace the process' memory space with a new program



دی ۹۹ / نویم / ۲۴

Process creation with C

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

POSIX

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
                      "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
                      NULL, /* don't inherit process handle */
                      NULL, /* don't inherit thread handle */
                      FALSE, /* disable handle inheritance */
                      0, /* no creation flags */
                      NULL, /* use parent's environment block */
                      NULL, /* use parent's existing directory */
                      &si,
                      &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

Windows

Process termination

➤ Child → Parent

○ Process' resources are deallocated when:

- *exit(n)*
- *return()* in *main()*

○ Catch exit status → *wait()*

- `pid = wait(&status);`

➤ Parent → Child

○ *abort()* → *ابرتو* (پس از که child پر کیل)

○ Why?

- Child has exceeded allocated resources
- Task assigned to child is no longer required
- The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Problems of process termination

➤ **zombie process**

- No parent waiting → child : terminated *orphaned Child*

➤ **orphan process**

- Parent termination without wait

➤ **Multi process example: Chrome Browser**

- Browser, Renderer, Plugins, etc



Interprocess communication (*IPC*)

Process:

- o independent vs. cooperating

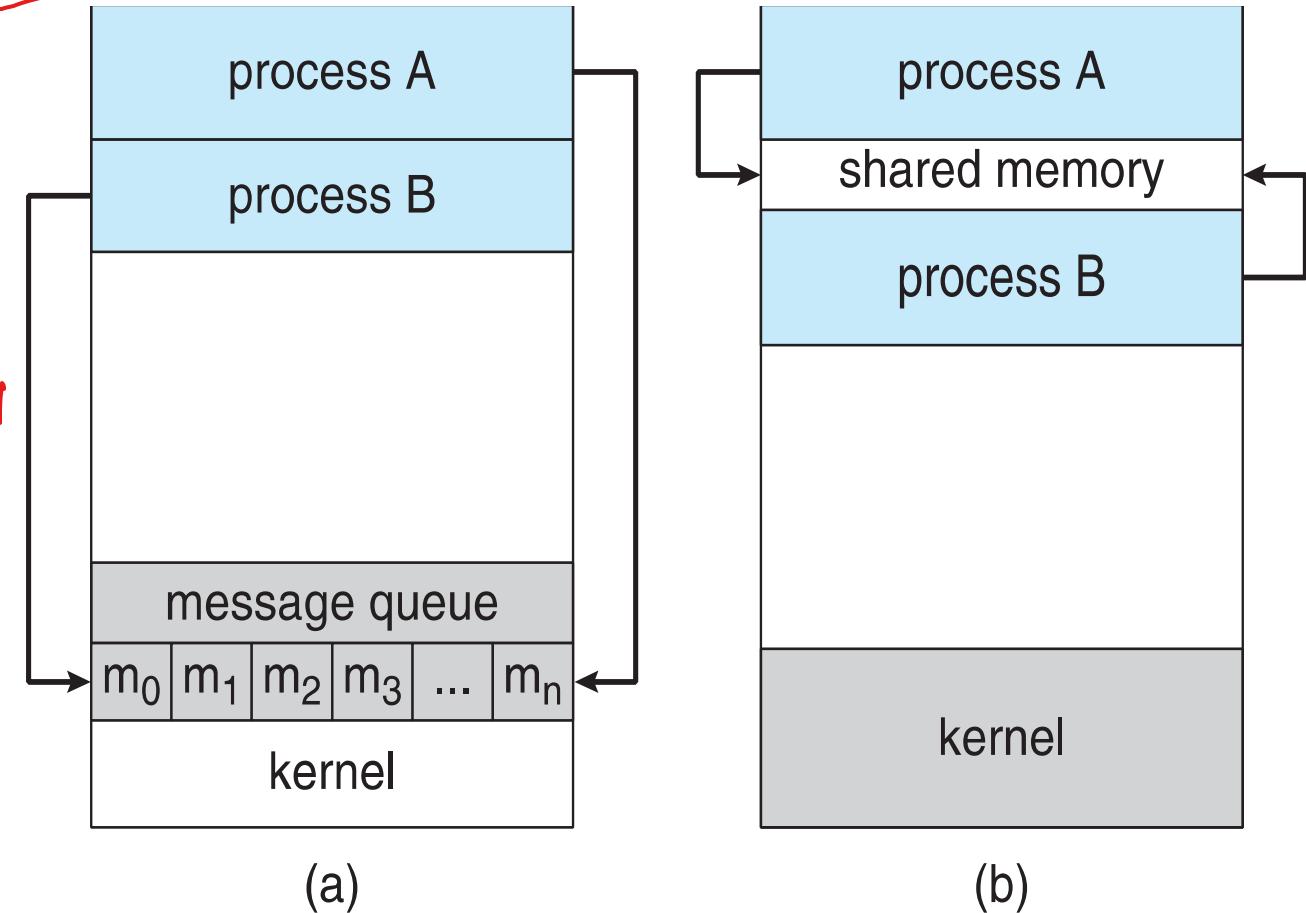
~~share data~~
 - objects, files, databases, ...

info sharing
 computation speedup (task < subtask)
 modularity
 convenience (fun)

Cooperating process:

- o Shared memory
- o Message passing

buffers
 bounded
 unbounded
 direct - unidirectional
 indirect - unidirectional
 bidirectional
 using mailbox (ports)



Circular buffer & producer-consumer problem

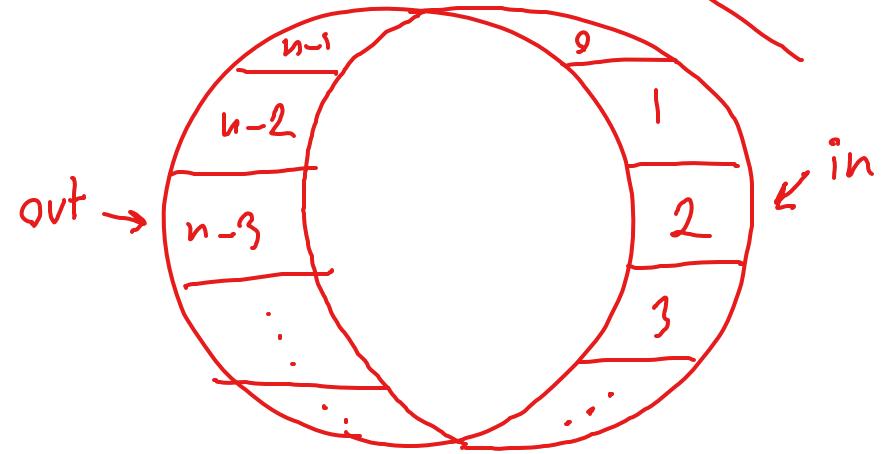
~~shared
mem.~~

in = out → empty buffer

*bounded
buffer*

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```



```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

```
item next_consumed;
while (true) {
    while (in == out) ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}
```

Message passing

➤ Direct communication (unidirectional)

- **send** (P , message) – send a message to process P
- **receive**(Q , message) – receive a message from process Q

➤ Indirect communication (uni & bidirectional)

- Messages are directed and received from mailboxes (or ports)
- Can be used by multiple processes
- Primitives are defined as:
 - **send(A, message)** – send a message to **mailbox A**
 - **receive(A, message)** – receive a message from **mailbox A**

Synchronization

- **Blocking vs. non-blocking**
- **Blocking is considered synchronous** → *نحوه نایم خود را خط بس!*
 - Blocking send
 - Blocking receive
- **Non-blocking is considered asynchronous** → *نحوه نایم null:*
 - Non-blocking send
 - Non-blocking receive
 - The receiver receives
 - A valid message
 - Null message
- **Different combinations possible**
 - If both send and receive are blocking, we have a rendezvous *مرا ملاقات*

POSIX examples of shared memory: (sender->receiver)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
/* the size (in bytes) of shared memory object */
const int SIZE = 4096;
/* name of the shared memory object */
const char *name = "OS"; *
/* strings written to shared memory */
const char *message_0 = "Hello";
const char *message_1 = "World!";

/* shared memory file descriptor */
int shm_fd;
/* pointer to shared memory obect */
void *ptr;

/* create the shared memory object */
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

/* configure the size of the shared memory object */
ftruncate(shm_fd, SIZE);

/* memory map the shared memory object */
ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

/* write to the shared memory object */
sprintf(ptr,"%s",message_0);
ptr += strlen(message_0);
sprintf(ptr,"%s",message_1);
ptr += strlen(message_1);

return 0;
}
```



```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
/* the size (in bytes) of shared memory object */
const int SIZE = 4096;
/* name of the shared memory object */
const char *name = "OS"; *
/* shared memory file descriptor */
int shm_fd;
/* pointer to shared memory obect */
void *ptr;

/* open the shared memory object */
shm_fd = shm_open(name, O_RDONLY, 0666);

/* memory map the shared memory object */
ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

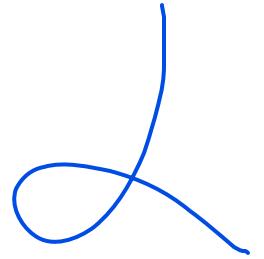
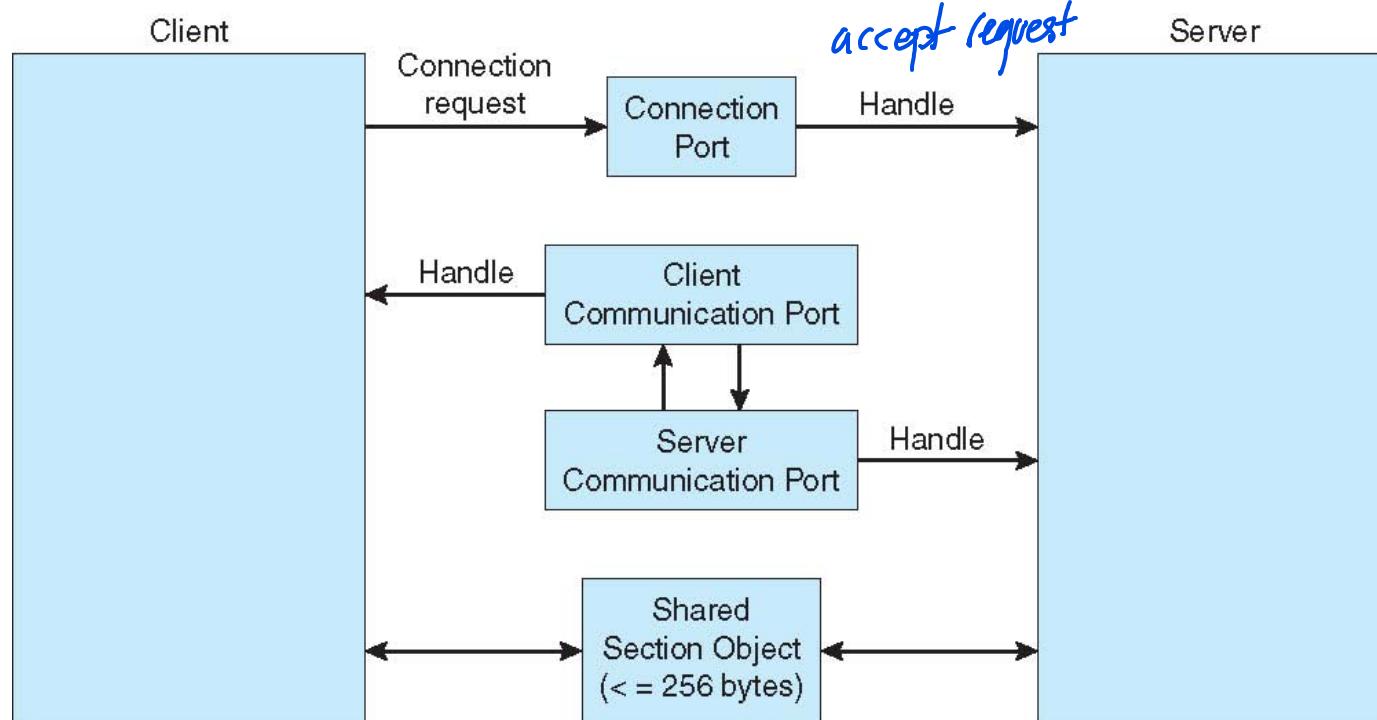
/* read from the shared memory object */
printf("%s",(char *)ptr);

/* remove the shared memory object */
shm_unlink(name);

return 0;
}
```

Local procedure calls in Windows (ALPC)

Advanced



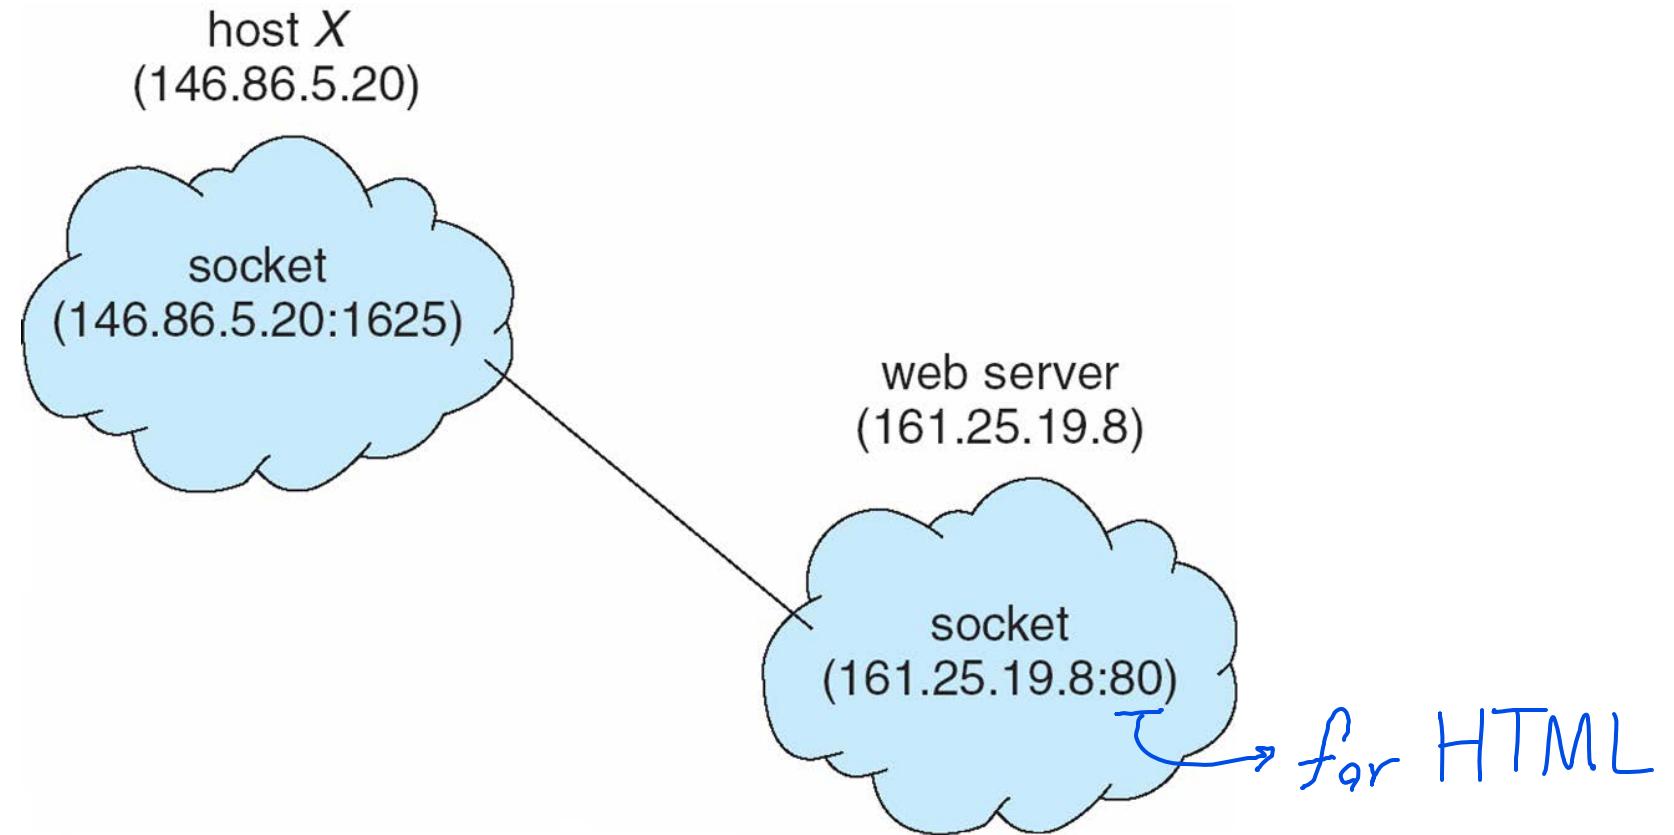
Communications in client-server systems

- Sockets
- Remote Procedure Calls (windows)
- Pipes
- Remote Method Invocation (Java)

Sockets

- A **socket** is defined as an endpoint for communication
- Concatenation of IP address and **port** – a number included at start of message packet to differentiate network services on a host
 - The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
 - Shouldn't be reserved*
 - Communication consists between a pair of sockets
 - All ports below 1024 are **well known**, used for standard services 
 - Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running

Socket communication



Sockets in Java

- Three types of sockets
 - Connection-oriented (TCP)
 - Connectionless (UDP)
 - MulticastSocket class—
data can be sent to multiple
recipients

- Consider this “Date” server:

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

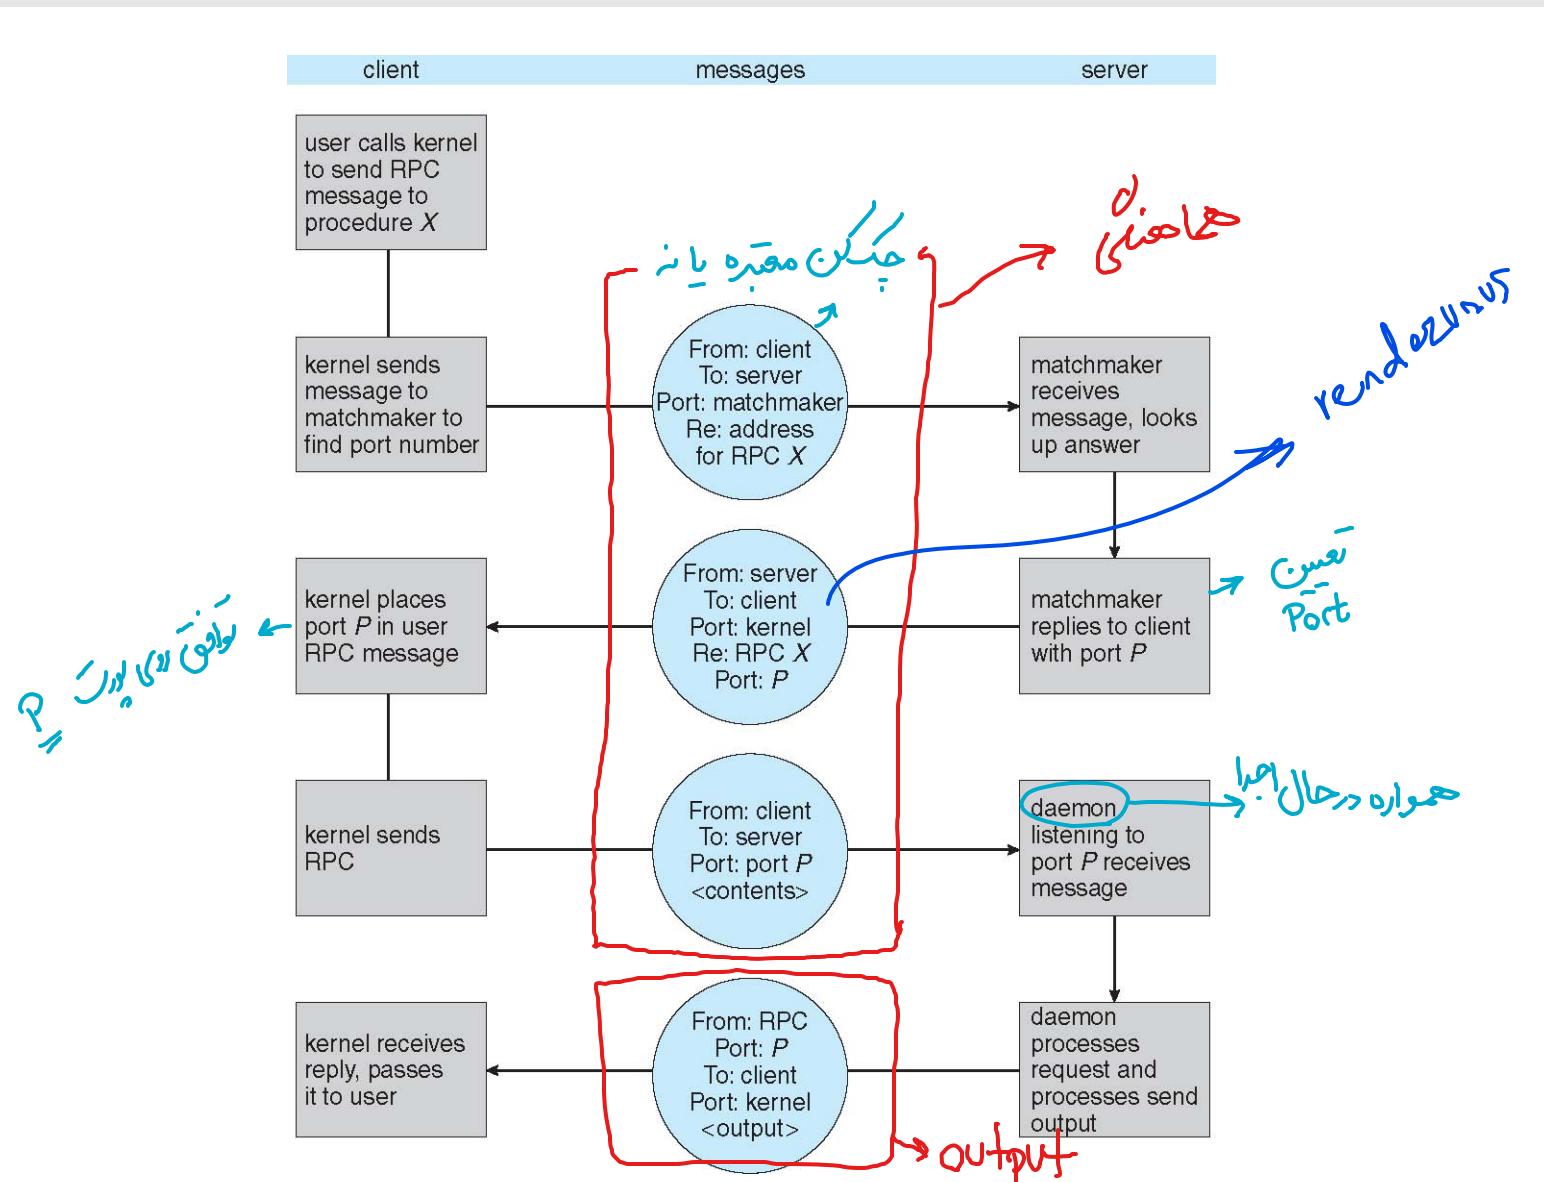
            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

Execution of RPC (Remote Procedure Call)



Pipes ($\#Parent, \#Child$)

➤ Acts as a conduit allowing two processes to communicate

➤ Issues:

- Is communication unidirectional or bidirectional?
- In the case of two-way communication, is it half or full-duplex?
- Must there exist a relationship (i.e., *parent-child*) between the communicating processes?
- Can the pipes be used over a network?

بَابِ حُرْنَانْ دُوْلَفِنْ

➤ Ordinary pipes (unidirectional)

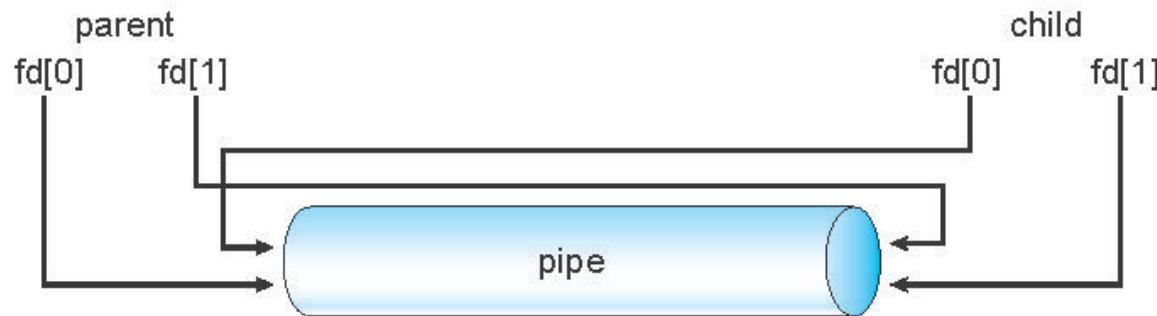
- cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.

➤ Named pipes (bidirectional)

- can be accessed without a parent-child relationship.

Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the write-end of the pipe)
- Consumer reads from the other end (the read-end of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes



- Windows calls these anonymous pipes
- See Unix and Windows code samples in textbook

Ordinary pipe (POSIX), parent-child

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>

#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1

int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;
    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    if (pid > 0) { /* parent process */
        /* close the unused end of the pipe */
        close(fd[READ_END]);
        /* write to the pipe */
        write(fd[WRITE_END], write_msg, strlen(write_msg)+1);
        /* close the write end of the pipe */
        close(fd[WRITE_END]);
    }
    else { /* child process */
        /* close the unused end of the pipe */
        close(fd[WRITE_END]);
        /* read from the pipe */
        read(fd[READ_END], read_msg, BUFFER_SIZE);
        printf("read %s", read_msg);
        /* close the write end of the pipe */
        close(fd[READ_END]);
    }
    return 0;
}
```

Ordinary pipe (windows), parent

```

#include <stdio.h>
#include <stdlib.h>
#include <windows.h>

#define BUFFER_SIZE 25

int main(VOID)
{
    HANDLE ReadHandle, WriteHandle;
    STARTUPINFO si;
    PROCESS_INFORMATION pi;
    char message[BUFFER_SIZE] = "Greetings";
    DWORD written;

    /* set up security attributes allowing pipes to be inherited */
    SECURITY_ATTRIBUTES sa = {sizeof(SECURITY_ATTRIBUTES),NULL,TRUE};
    /* allocate memory */
    ZeroMemory(&pi, sizeof(pi));

    /* create the pipe */
    if (!CreatePipe(&ReadHandle, &WriteHandle, &sa, 0)) {
        fprintf(stderr, "Create Pipe Failed");
        return 1;
    }

    /* establish the STARTINFO structure for the child process */
    GetStartupInfo(&si);
    si.hStdOutput = GetStdHandle(STD_OUTPUT_HANDLE);

    /* redirect standard input to the read end of the pipe */
    si.hStdInput = ReadHandle;
    si.dwFlags = STARTF_USESTDHANDLES;

    /* don't allow the child to inherit the write end of pipe */
    SetHandleInformation(WriteHandle, HANDLE_FLAG_INHERIT, 0);

    /* create the child process */
    CreateProcess(NULL, "child.exe", NULL, NULL,
        TRUE, /* inherit handles */
        0, NULL, NULL, &si, &pi);

    /* close the unused end of the pipe */
    CloseHandle(ReadHandle);

    /* the parent writes to the pipe */
    if (!WriteFile(WriteHandle, message,BUFFER_SIZE,&written,NULL))
        fprintf(stderr, "Error writing to pipe.");

    /* close the write end of the pipe */
    CloseHandle(WriteHandle);

    /* wait for the child to exit */
    WaitForSingleObject(pi.hProcess, INFINITE);
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
    return 0;
}

```

Ordinary pipe (windows), child

```
#include <stdio.h>
#include <windows.h>

#define BUFFER_SIZE 25

int main(VOID)
{
    HANDLE Readhandle;
    CHAR buffer[BUFFER_SIZE];
    DWORD read;

    /* get the read handle of the pipe */
    ReadHandle = GetStdHandle(STD_INPUT_HANDLE);

    /* the child reads from the pipe */
    if (ReadFile(ReadHandle, buffer, BUFFER_SIZE, &read, NULL))
        printf("child read %s",buffer);
    else
        fprintf(stderr, "Error reading from pipe");

    return 0;
}
```

Named pipes

- Named Pipes are more powerful than ordinary pipes (?) → *پایه ایزبی رفن* (Parent) *کلیو ایزبی* (Child)
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems

Questions?

