



Main Memory Management

Hamid R. Zarandi
h_zarandi@aut.ac.ir

Why memory management?

- CPU utilization and increase of performance
- → need of scheduler to schedule multiple processes
- → they should be resident in RAM (physical memory)

- Why need of memory management?
 - Having multiple processes
 - Protection (processes from one another, processes from OS)

- Protection should be done by Hardware than OS (why?) \Rightarrow
 - Performance penalty and latency

How to protect process memory space?

➤ Each process has a **separate memory space**

➤ To protect processes' spaces

- Determining **legal address**

- **Base register** (پایه) : smallest legal physical memory address
- **Limit register** (حد) : size of the range

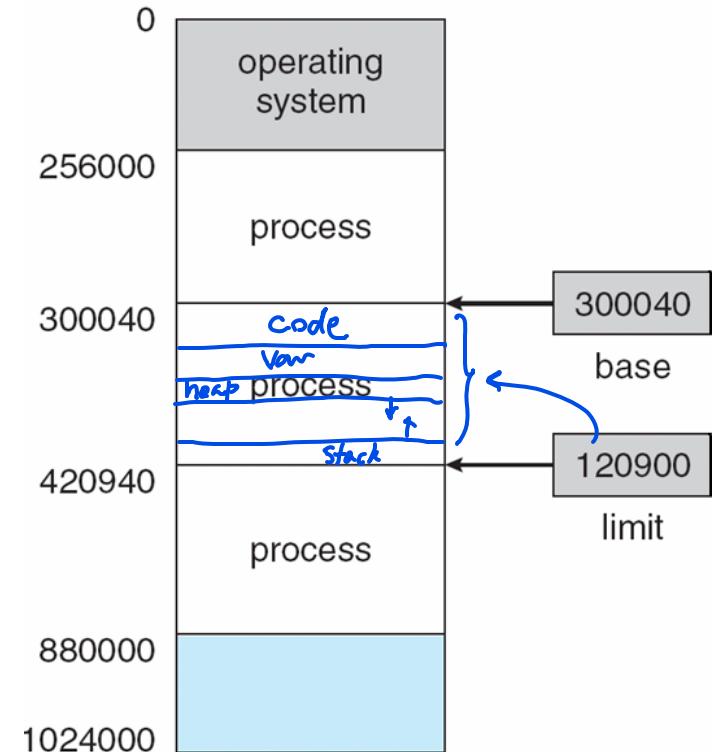
- Example:

- Base register = 300040
- Limit register = 120900

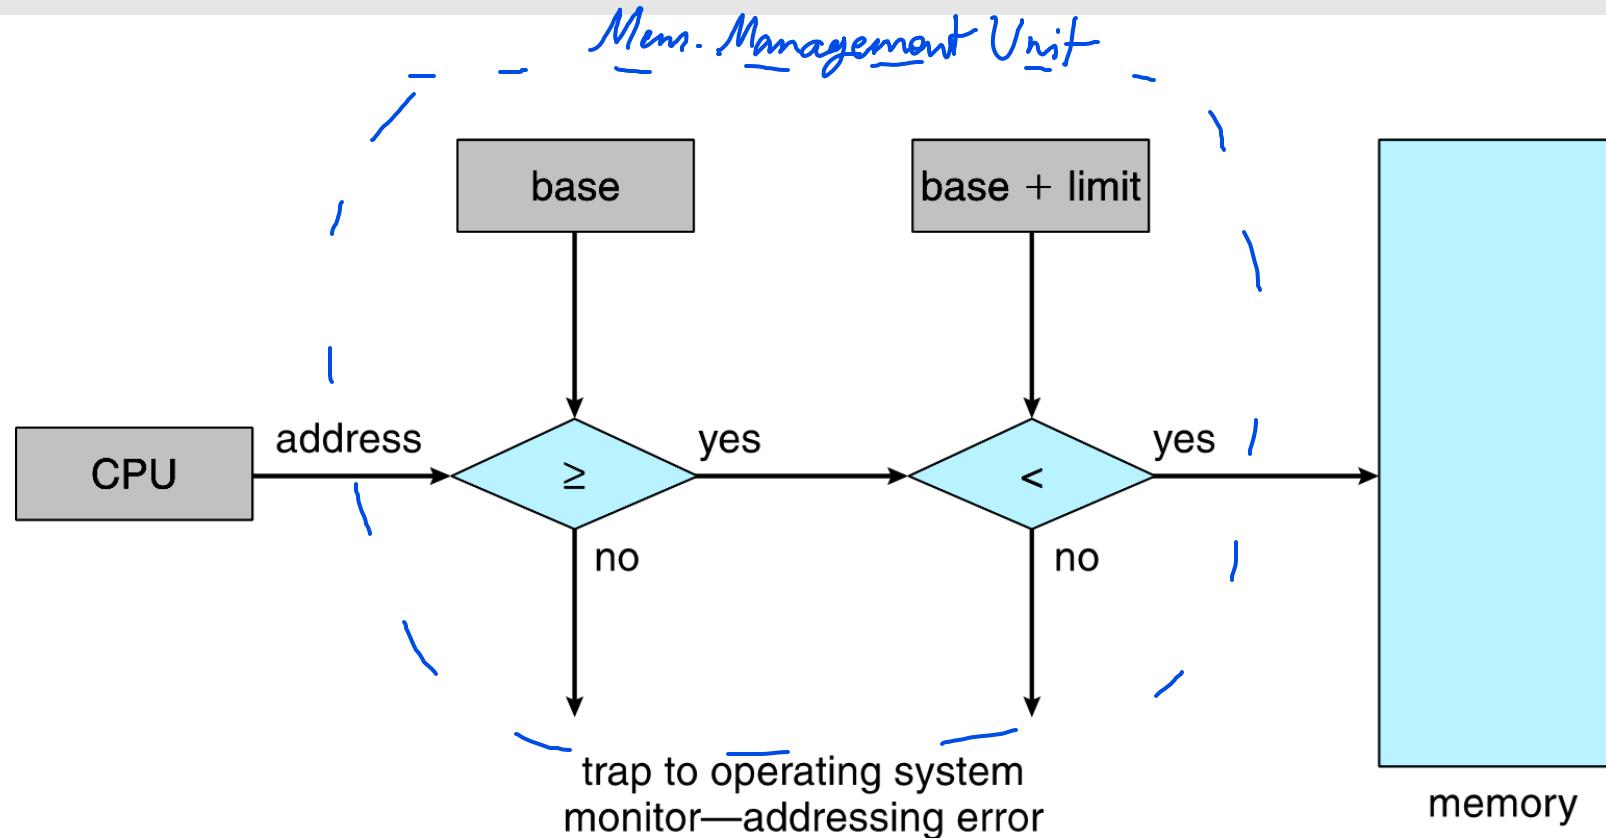
- **Legal address space:**

$$\text{Base} \leq (\text{any address}) < \text{Base} + \text{Limit}$$

- Can easily be checked by hardware



HW address protection (base & limit regs.)



Any illegal address generates a **trap exception** which is known as **fatal error**

Who loads Base and Limit registers?

→ OS

Address binding

Address binding

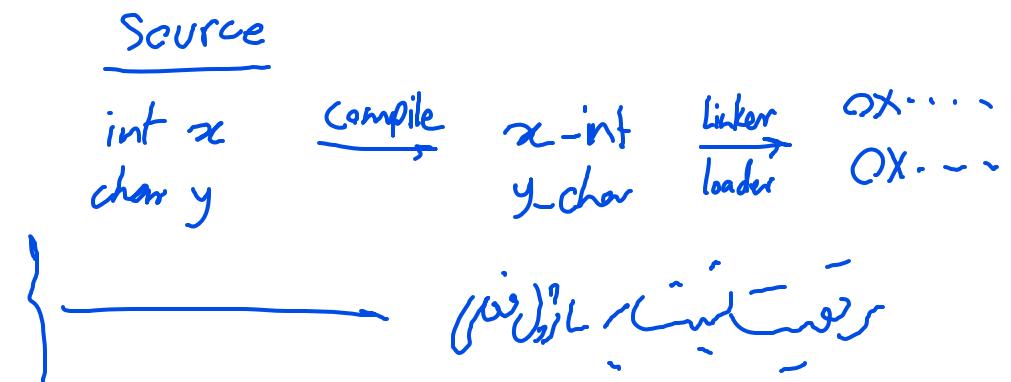
physical → MM^(S)
 logical → (CPV) ^{physical} → input queue (in Ready queue)

➤ Input Queue

- Processes on disk that are waiting to be brought into memory during execution (Part of ready queue which is on disk)

➤ How to put a process in a physical address?

- Addresses in source program are symbolic
 - Example: *count* variable
- A compiler binds them to relocatable addresses
 - Example: 14 bytes from beginning of this module
- Linker and loader bind them next to absolute addresses
 - Example: 74014



How to bind inst./data to mem. address?

➤ Compile time

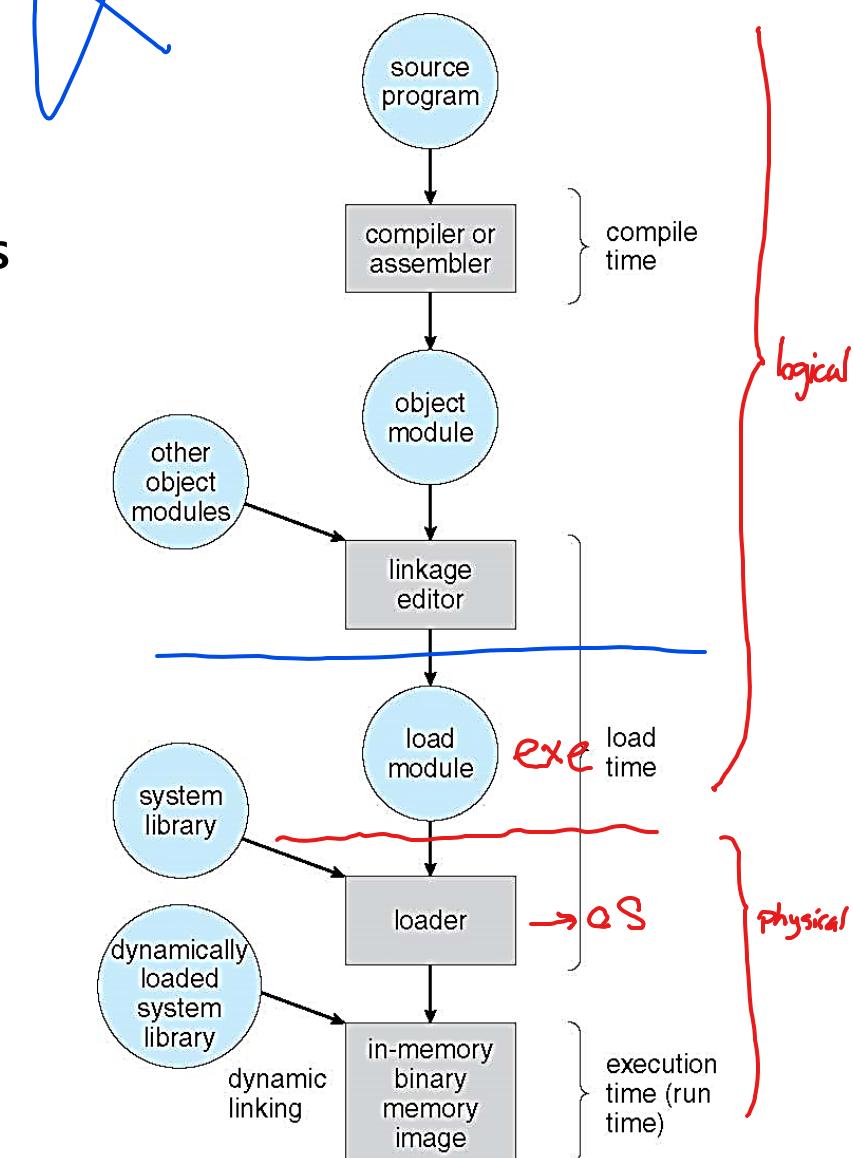
- If memory location known *a priori*, **absolute code can be generated**; must recompile code if starting location changes
 - Example: **COM files** in MS-DOS

➤ Load time

- Must generate **relocatable code** if memory location is not known at compile time

➤ Execution time

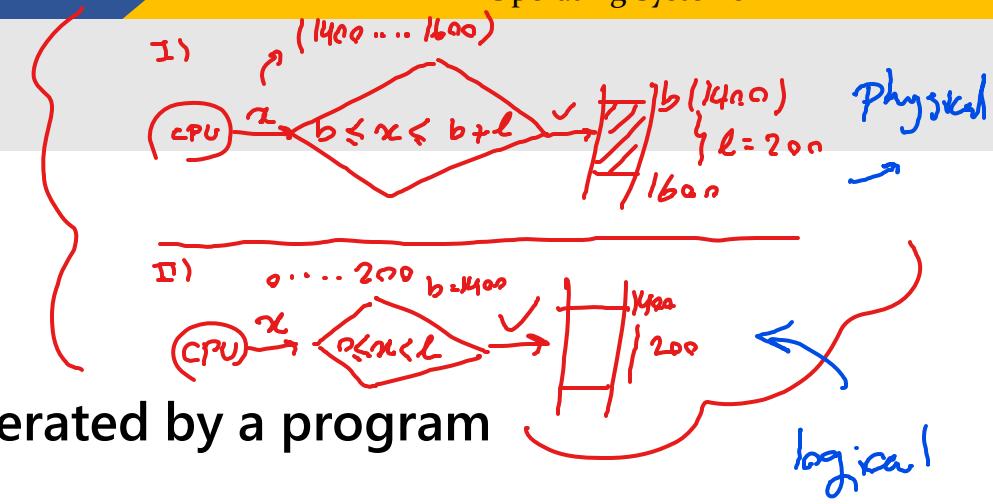
- Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - Need **hardware support for address maps**
 - ✓ Example: **Base-Limit registers**



Logical vs. physical address space

➤ Logical address (CPU address)

- CPU logically sees addresses
- **Logical address space:** set of all **logical addresses** generated by a program

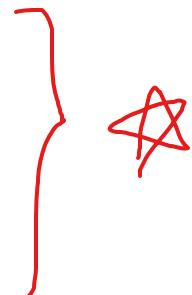


➤ Physical address (Memory address)

- Address of memory line
- **Physical address space:** set of all **physical addresses** generated by a program

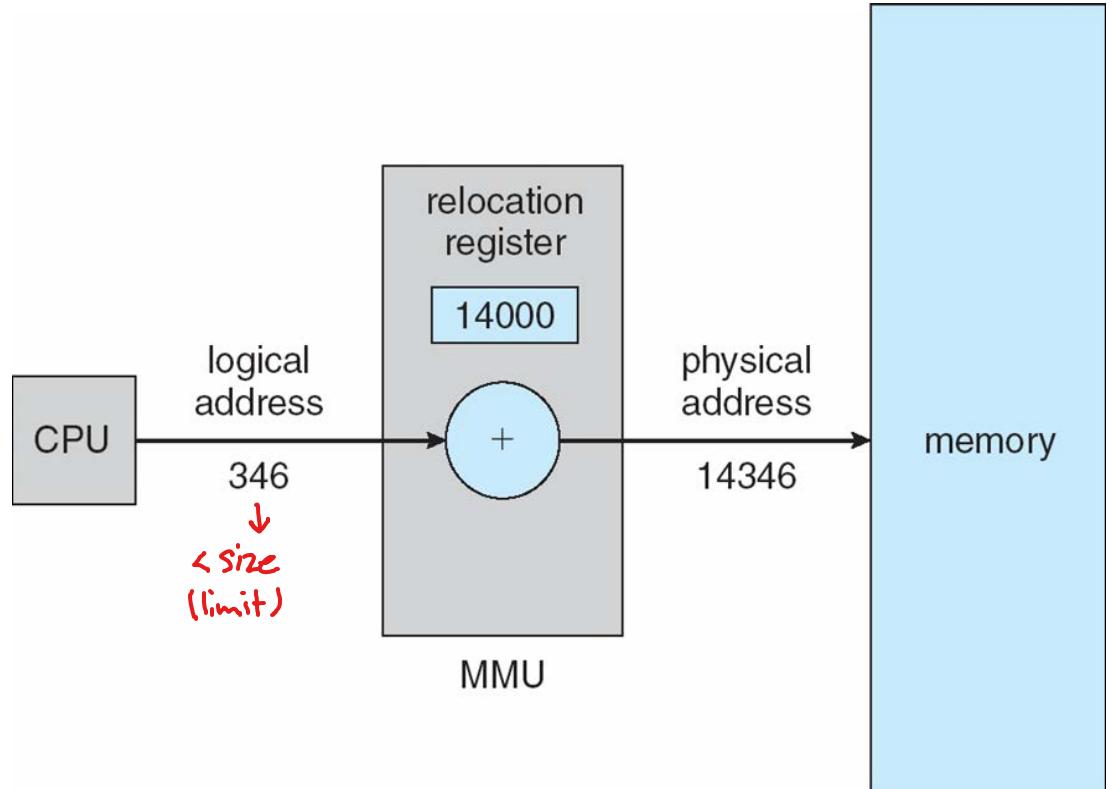
➤ May be equal or not

- **Equal:** **compile-time** and **load-time** address bindings
- **Not equal:** **execution-time** address binding
 - In this case, **logical address** is said **virtual address**



What is MMU (Memory Management Unit)?

- MMU is responsible of mapping **virtual (logical) address** to **physical address**
- Simple version: **Base-Limit registers**
- Here **Base register** is called **relocation register**



Dynamic linking and loading

Dynamic loading

- **Dynamic loading**
 - A routine is **not loaded** until it is **called**
 - All routines except **main()** are kept on disk
 - Better space utilization
 - Some routines are **infrequently** needed: **error functions** *((catch))*
 - In the case of call, if it is not loaded, linking loader first loads it to the memory and update program's address table to reflect this change

- is called *linked binary executable*
- kept on disk
- / needed: *error functions* \rightarrow *(catch)*
- loaded, linking loader first loads it to the address table to reflect this change

Dynamic linking

➤ Dynamic linking (Dynamically linked libraries)

- System libraries that are **linked** to user programs when they are **run**
- **No waste** of memory and disk space
- Example: **Language subroutines**
- **Stub**: small piece of code used to locate dynamic linked libraries
 - Replace itself with the address of the routine & executes the routine

- Also known as **Shared Libraries**

➤ Static linking

- System libraries are linked to user programs during **compile time**

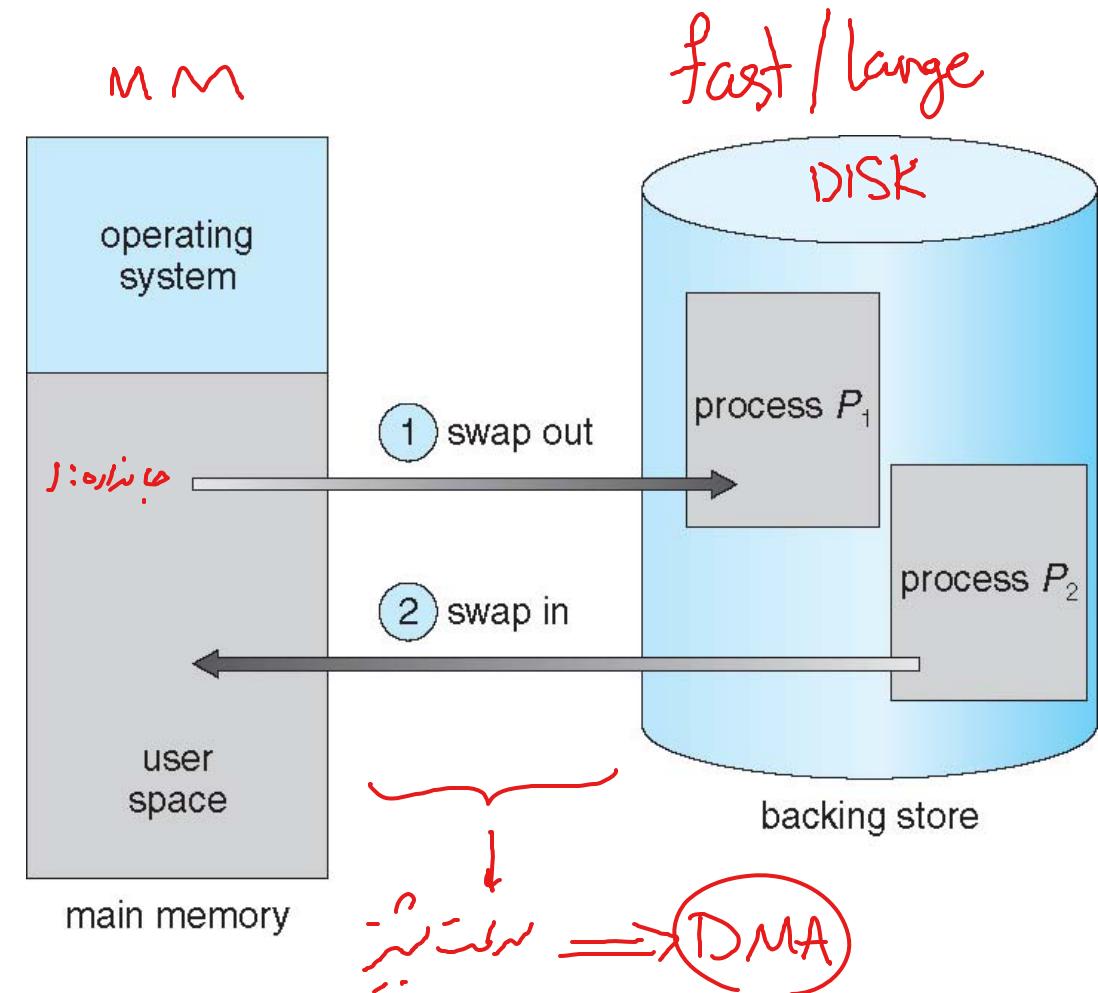
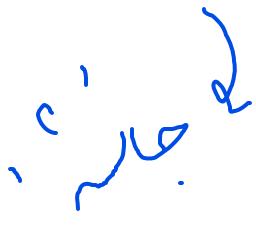
Swapping

Swapping

➤ **Swapping** is performed temporarily between memory and backing store

➤ **Backing store** – fast disk large enough to accommodate copies of **all** memory images for all users; must provide **direct** access to these memory images

➤ Possible of **total process memory spaces exceeds real physical memory**
 → Use of Backing store



Swapping issues

- Does the swapped out process need to swap back in to same physical addresses?
 - Depends on address binding method
 - Plus consider pending I/O to/from process memory space
- Pending IO: cannot swap out as IO would occur to wrong swapped in process!
 - Or always transfer IO to kernel space, then to IO device; this is known as **double buffering**; adds **overheads**

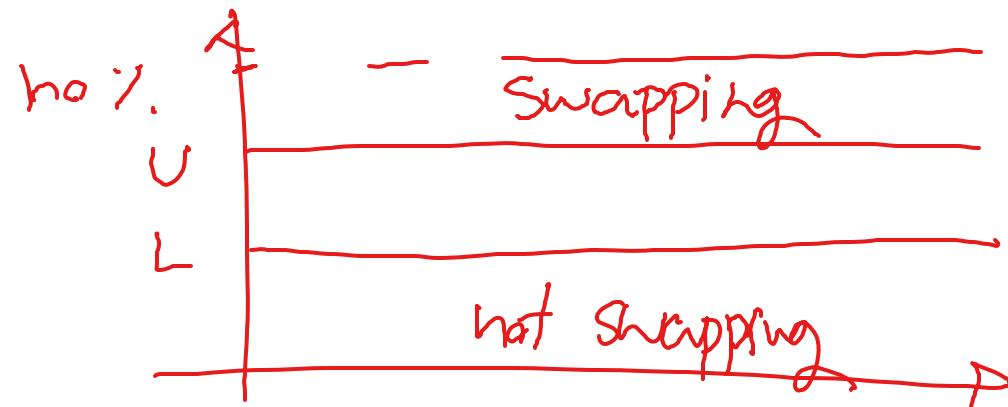
(ج) تبدیل مجدد شدن I/O

Swapping and modern OSs

➤ Standard swapping is not used in Linux & Windows!

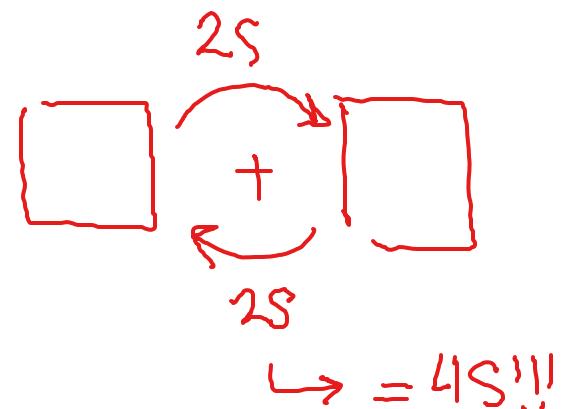
- Modified version is used:

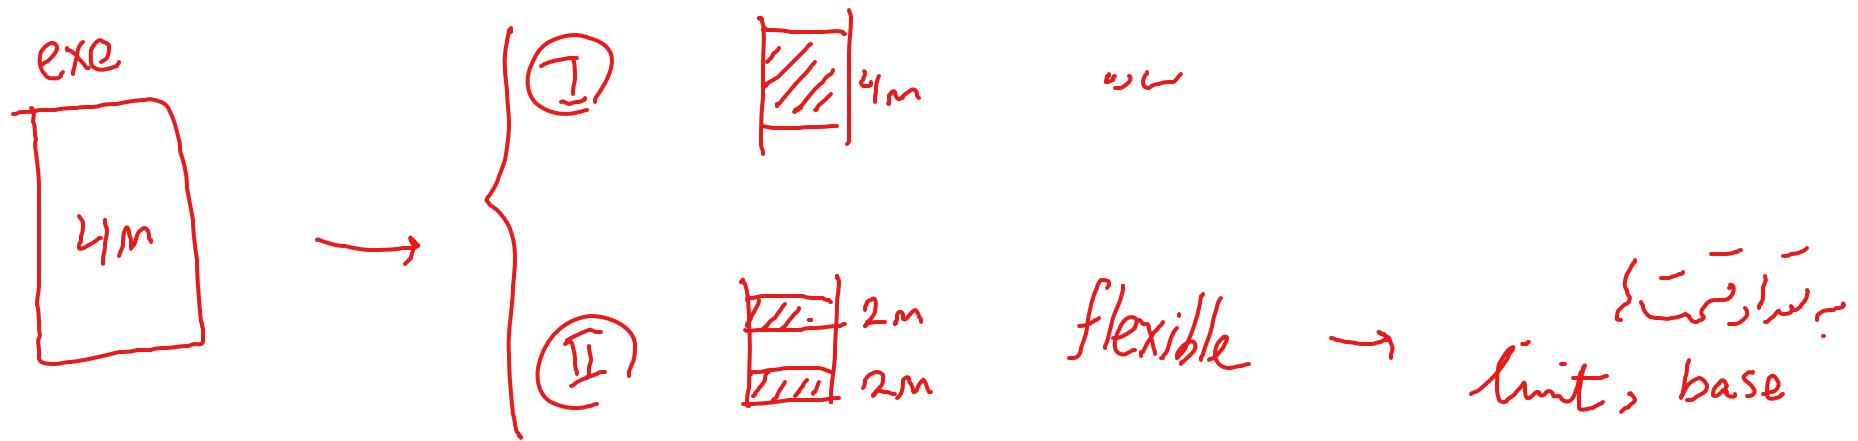
- Swap only when free memory is extremely low (less than threshold)
- Disable swapping when free memory is more than threshold



Swapping cost

- Major part of swap time is **transfer time**
- Total transfer time is proportional to the amount of memory swapped.
- Example:
 - 100 MB process swapping to hard disk with transfer rate of 50 MB/sec
 - Swap out time of 2(sec)+swap in of same size process
 - Total context switch swapping component time of 4 sec.





Memory allocation

Memory allocation

➤ Different types



1. Contiguous

- Each process is in a **single section** of memory that is **contiguous** to sections of others

مکانی متوالی
مکانی مجزا



2. Segmentation



- Each process is divided into different **segments**; each one is located in different part

3. Paging



- Each process is divided into **same-small-size pages**; **some** of them are swapped in/out

مکانی مجزا
مکانی متوالی

Criteria and problems

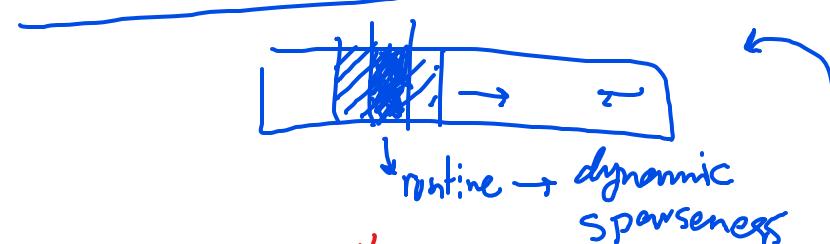
➤ CPU Utilization

- Utilization = Percent of time a CPU is busy = $\frac{\text{CPU time used}}{\text{Total time system is busy}}$

- Some times is $\frac{\text{CPU time used}}{\text{CPU time used} + 2 * \text{swapping time}}$
- Some times is $\frac{\text{CPU time used}}{\max(\text{CPU time used}, 2 * \text{swapping time})}$

➤ Fragmentation

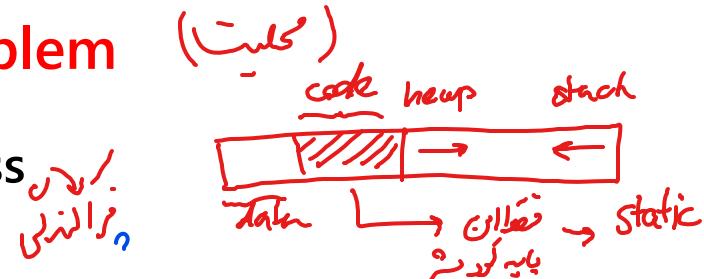
- External
- Internal



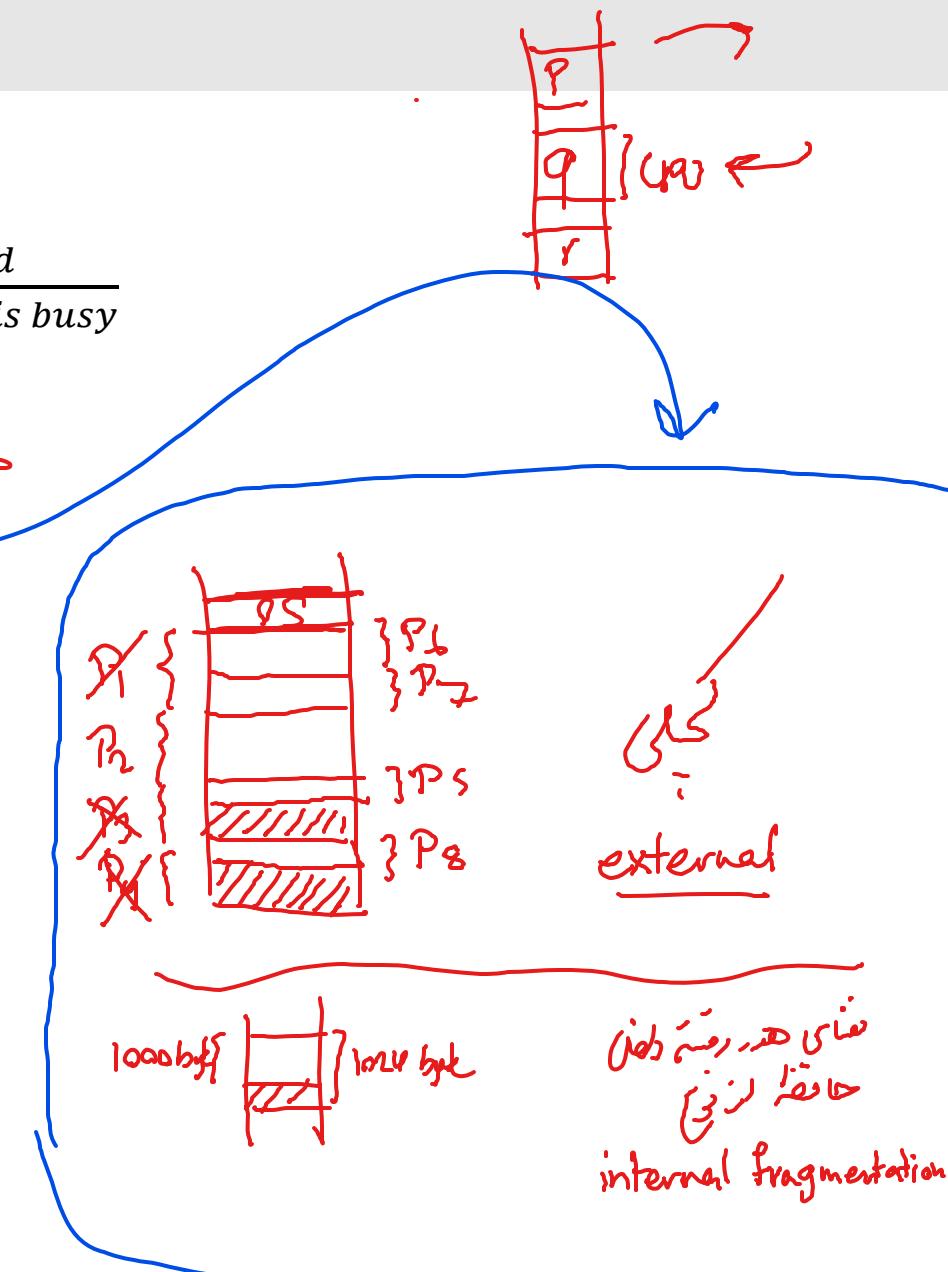
Time Sharing
Round Robin
MS to MS
CPU idle

➤ Process locality problem

- Static sparseness
- Dynamic sparseness



➤ Code/data sharing problem and protection problem

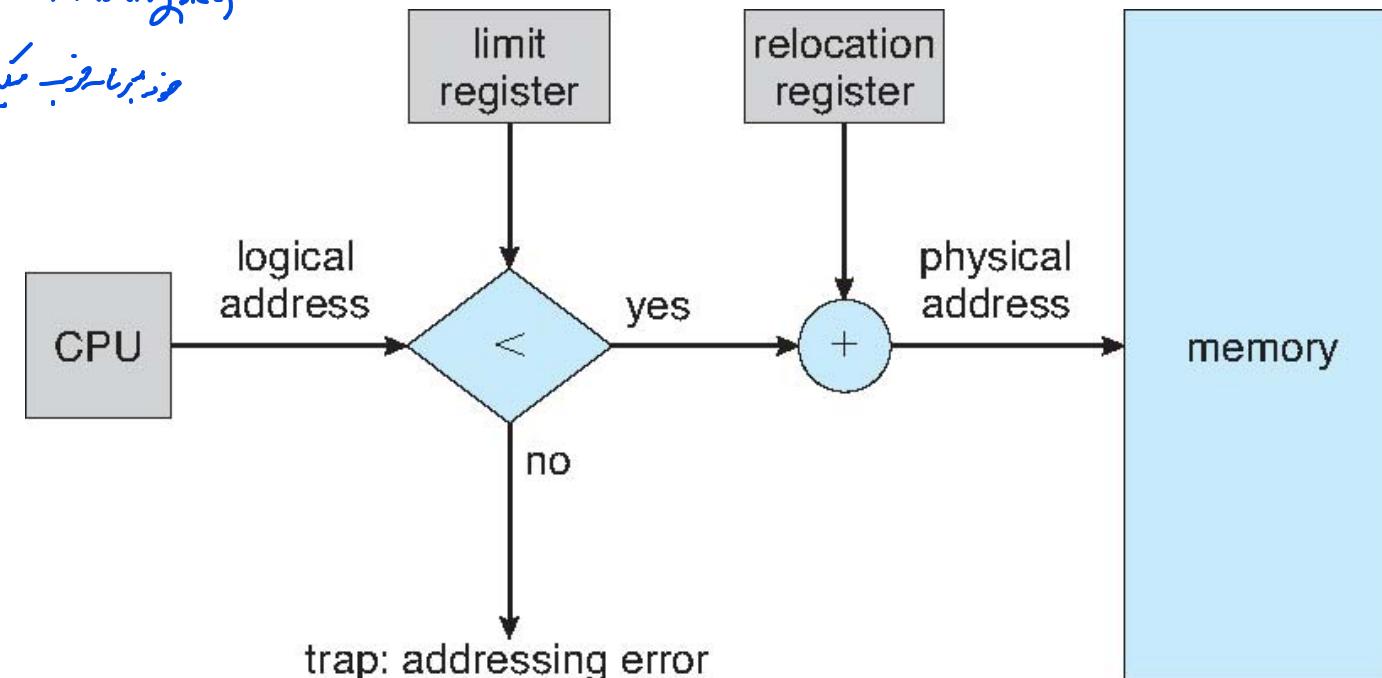
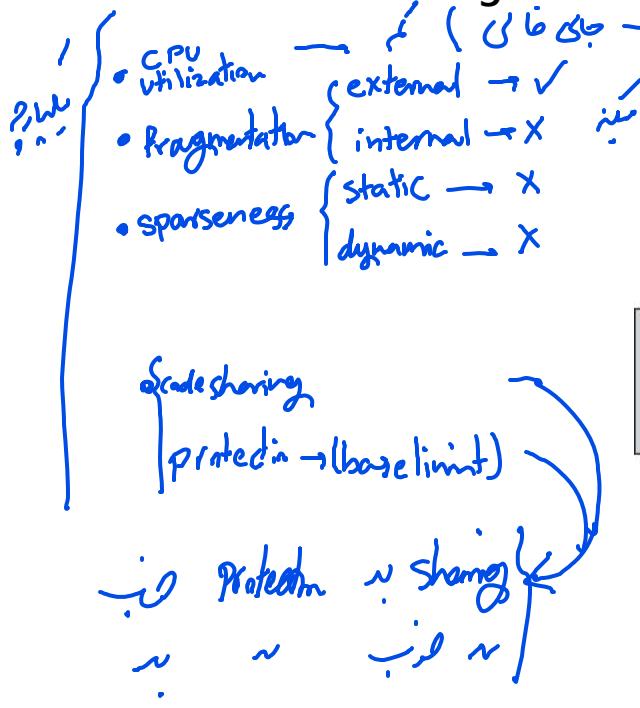


1) Contiguous allocation

➤ **Relocation registers** used to protect user processes from each other, and from changing operating-system code and data

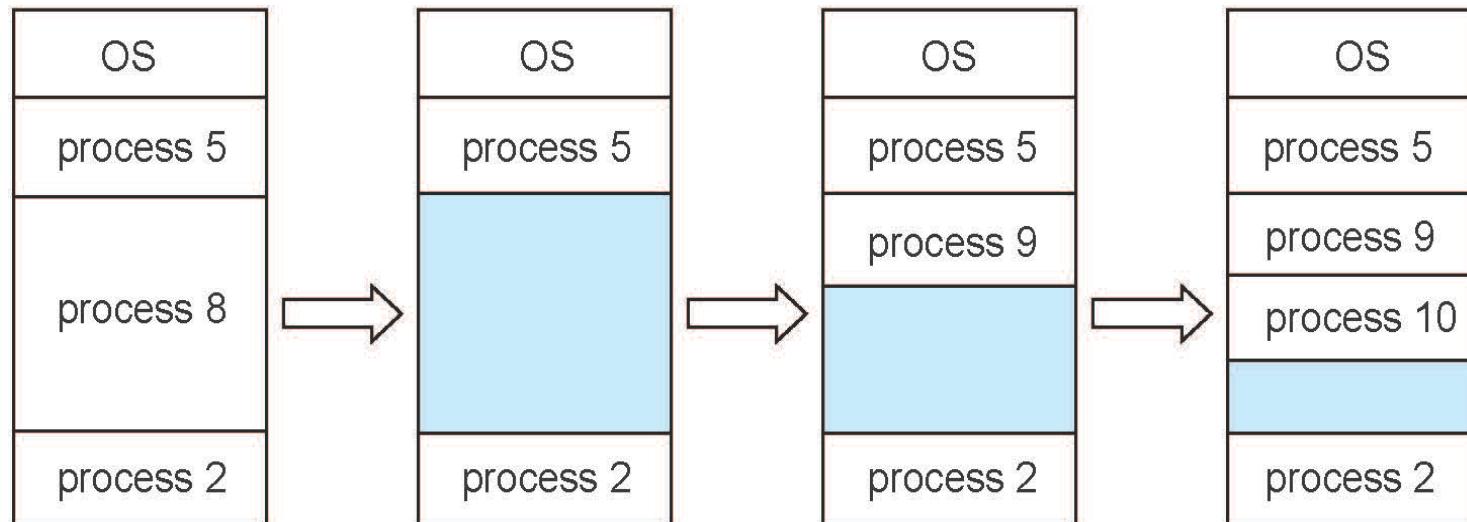
- Base register contains value of **smallest physical address**
- Limit register contains **range of logical addresses**

- Each logical address must be less than the limit register



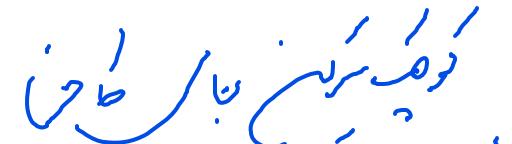
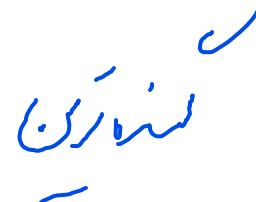
Multiple-partition memory allocation

- Degree of multiprogramming limited by number of partitions
- Variable-partition sizes for efficiency (sized to a given process' needs)
- Hole – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Operating system maintains information about:
 - a) allocated partitions b) free partitions (hole)



Dynamic storage-allocation problem

➤ How to satisfy a request of size n from a list of free holes?

- **First-fit:** Allocate the *first* hole that is big enough 
- **Best-fit:** Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole 
- **Worst-fit:** Allocate the *largest* hole; must also search entire list
 - Produces the largest leftover hole 

First-fit and best-fit better than worst-fit in terms of speed and storage utilization

Fragmentation problem

➤ External Fragmentation ✓

- Total memory space exists to satisfy a request, but it is not contiguous

➤ Internal Fragmentation ✗

- Not whole of allocated memory is filled by the process memory

➤ First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation

- $1/3$ may be unusable -> 50-percent rule

$$\frac{2}{3} \left(\underbrace{\text{left, min, } \alpha}_{\text{---}} \right) - X$$

$$\rightarrow \alpha = \frac{3}{2} X$$

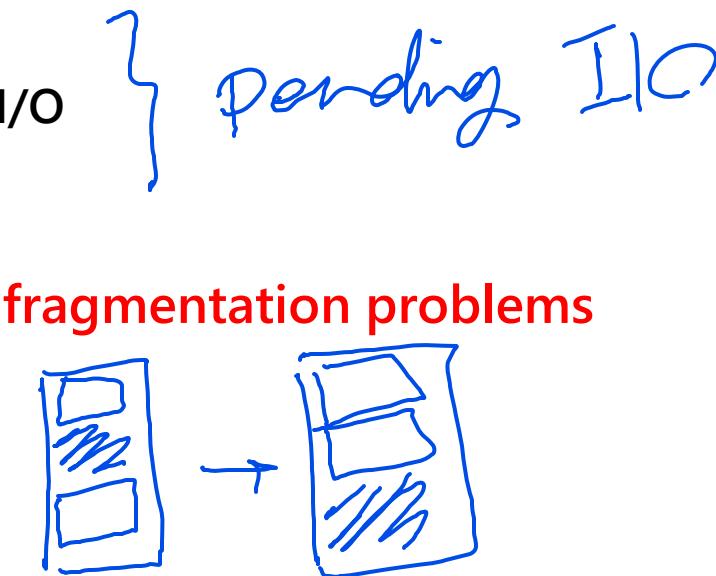
→ $\frac{1}{2} \times \frac{3}{2} X = \frac{3}{4} X$

✗

Fragmentation solution

➤ Compaction

- Shuffle memory contents to place all free memory together in one large block
- Compaction is possible *only* if relocation is **dynamic**, and is done at **execution time**
- I/O problem
 - Latch job in memory while it is involved in I/O
 - Do I/O only into OS buffers
- Now consider that backing store has same fragmentation problems



2) Segmentation

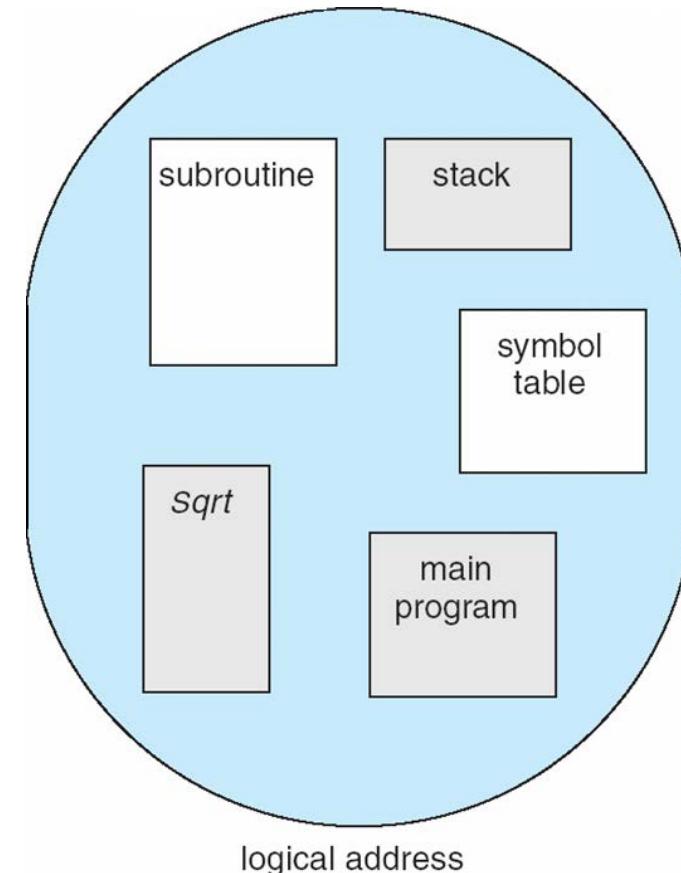
weis \leftarrow wiss

➤ Memory-management scheme that supports **user view of memory**

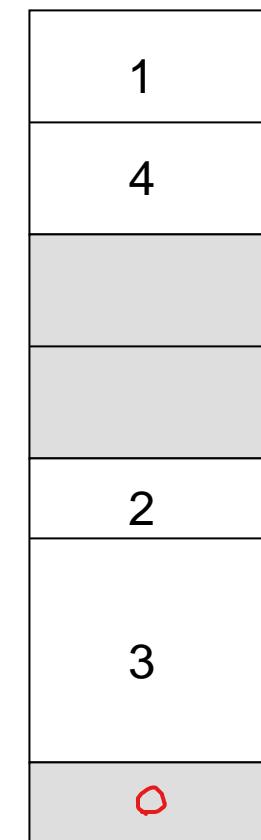
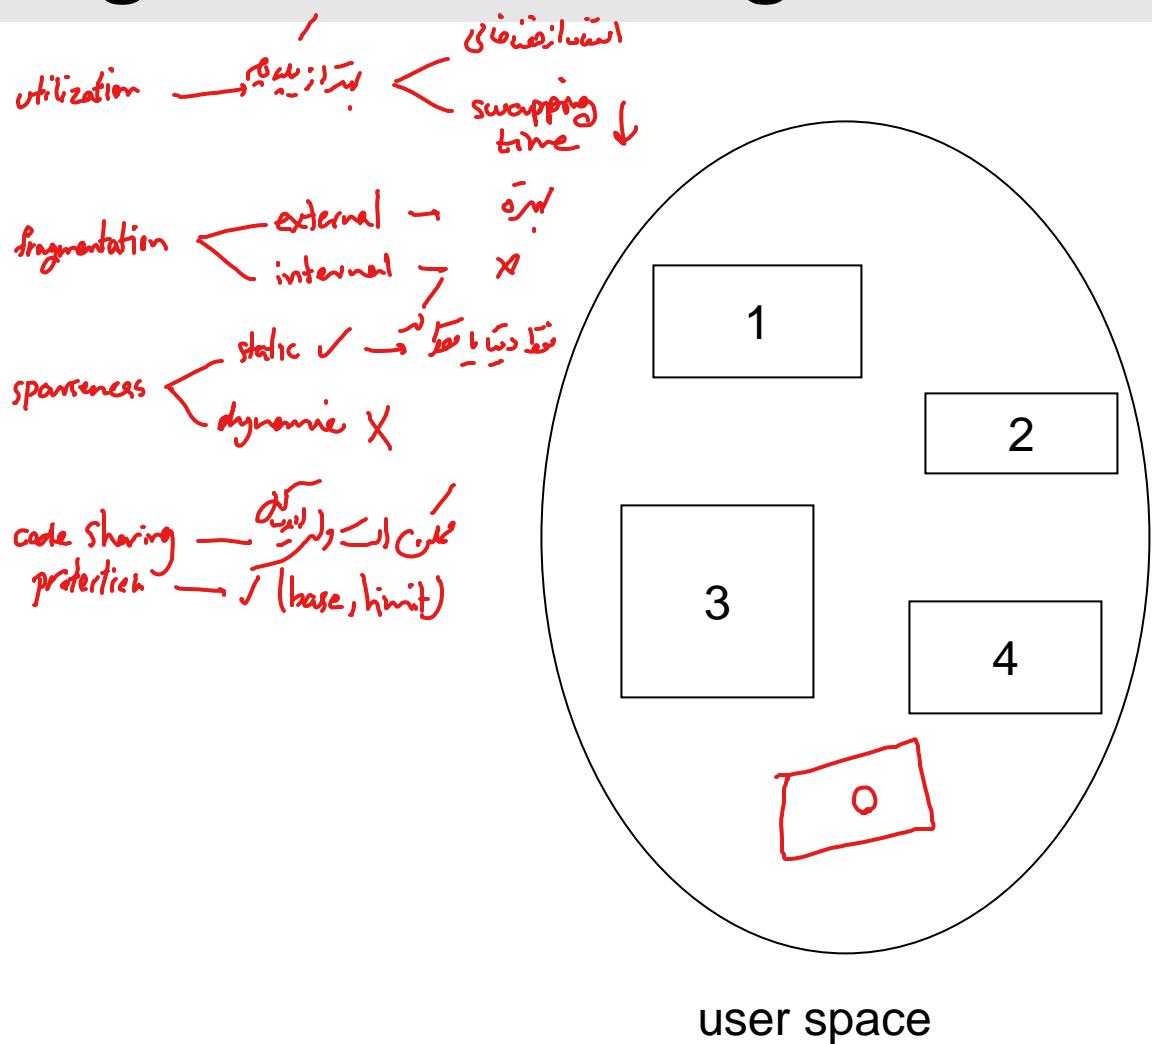
➤ A program is a collection of **segments**

○ A segment is a logical unit such as:

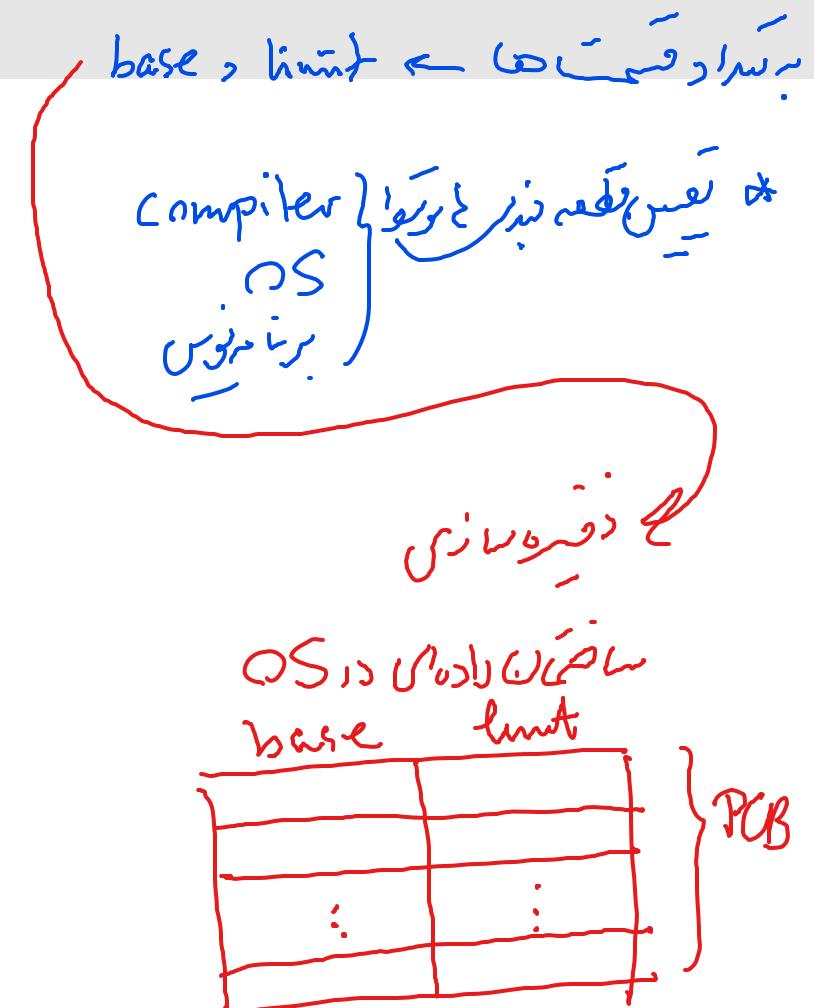
- main program
- procedure
- function
- method
- object
- local variables, global variables
- common block
- stack
- symbol table
- arrays



logical view of segmentation



physical memory space

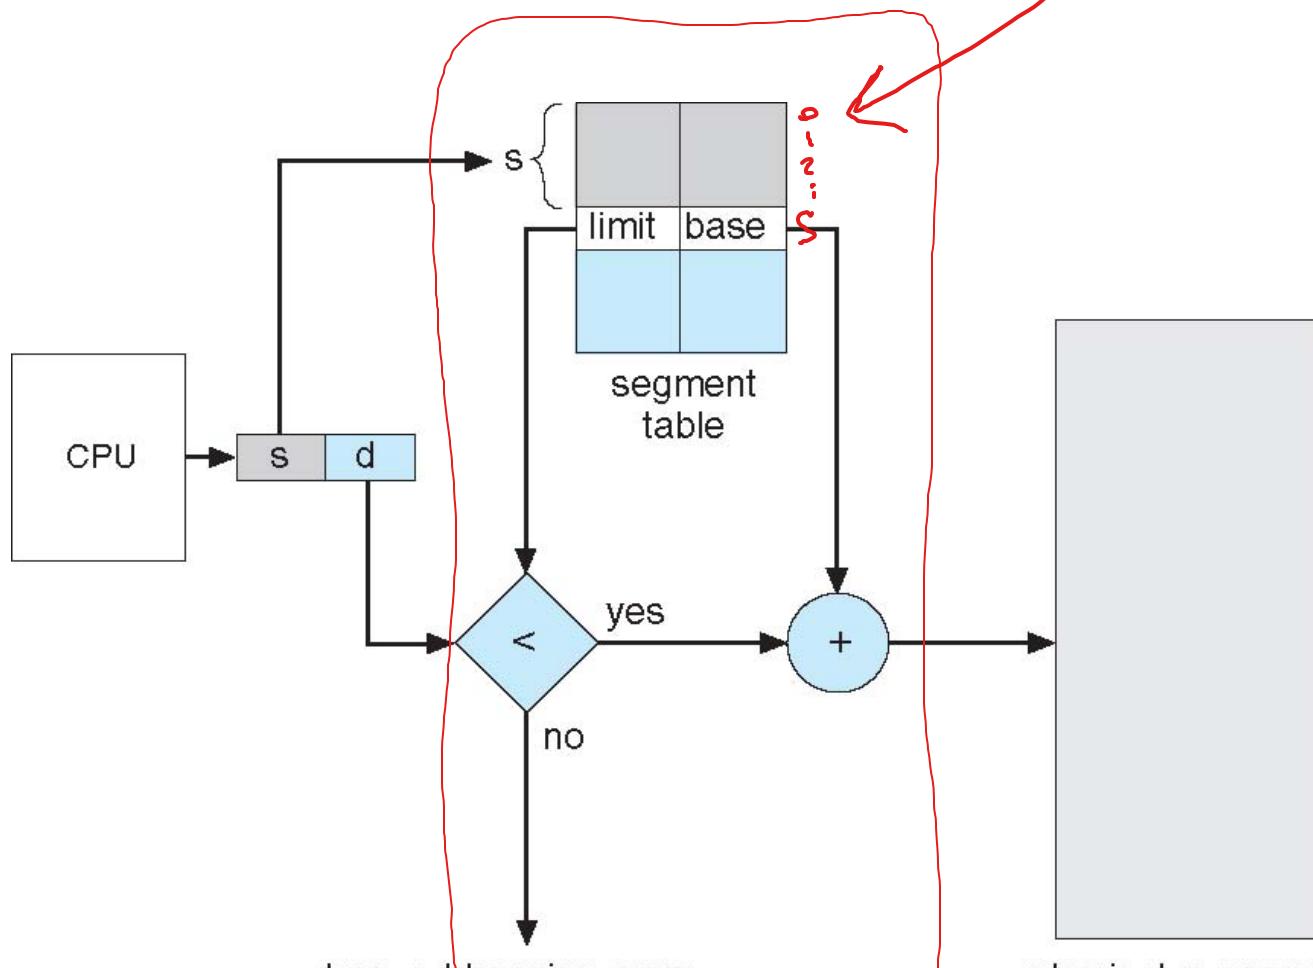


Segmentation implementation

- Logical address consists of a two tuple:
<segment-number, offset>

➤ Segment table

- Maps **two-dimensional physical addresses**; each table entry has:
 - **base** – contains the starting physical address where the segments reside in memory
 - **limit** – specifies the length of the segment

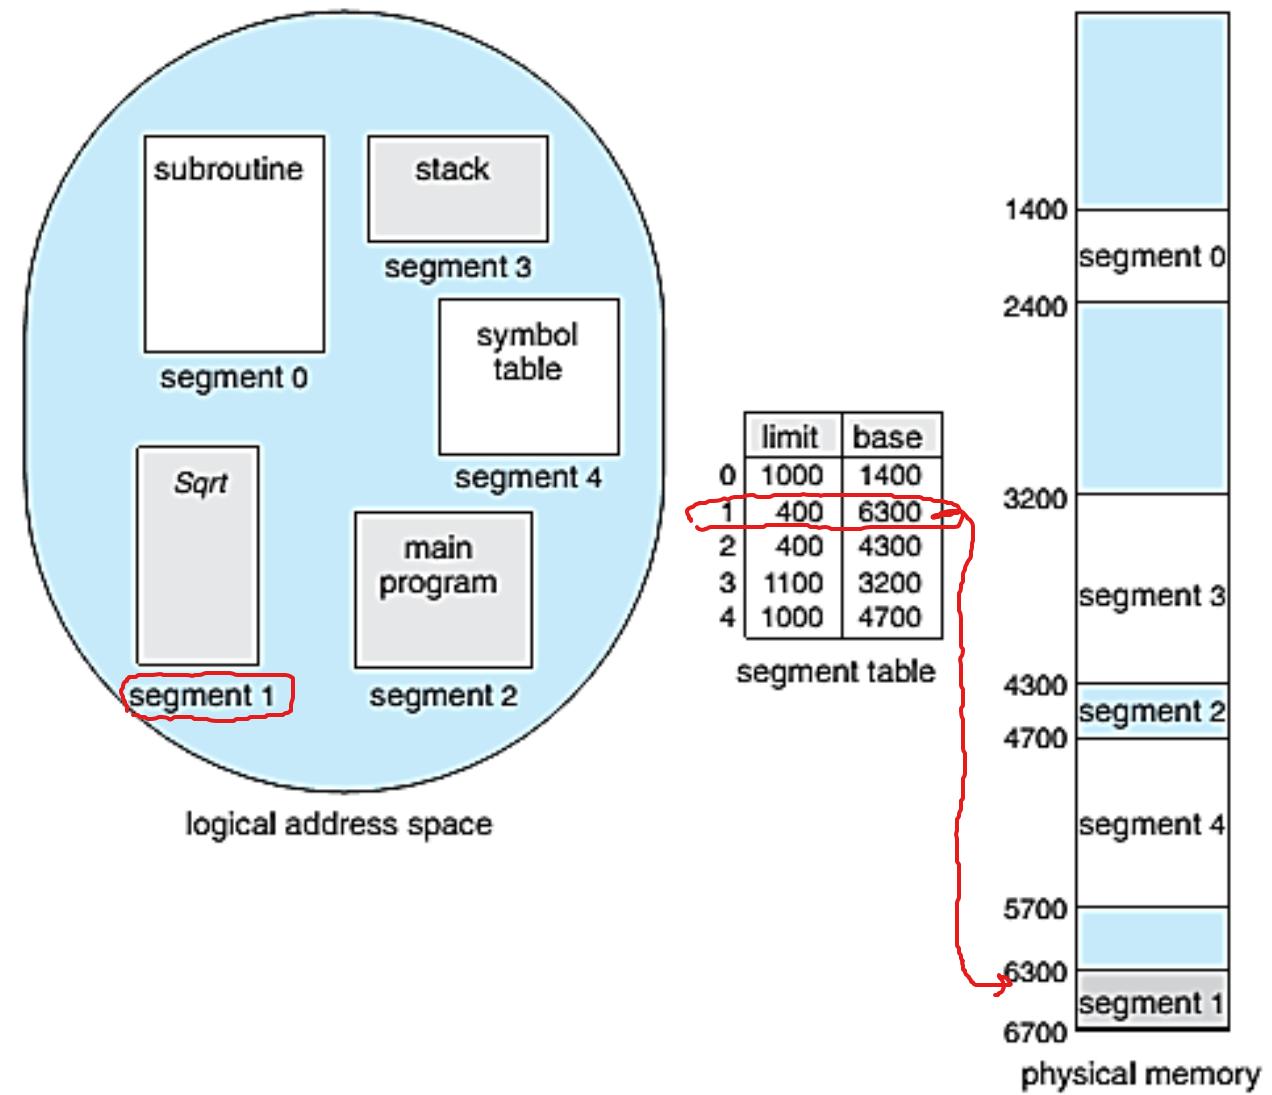


Example of segmentation

- A reference to byte 53 of segment 2:
 $4300+53=4353$

- A reference to byte 852 of segment 3:
 $3200+852=4052$

- A reference to byte 1222 of segment 0:
Trap to OS!



3) Paging

➤ Noncontiguous memory allocations:

- Segmentation
- Paging

➤ Paging avoids external fragmentation, and need of compaction, whereas segmentation does not.

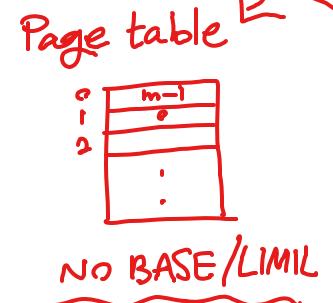
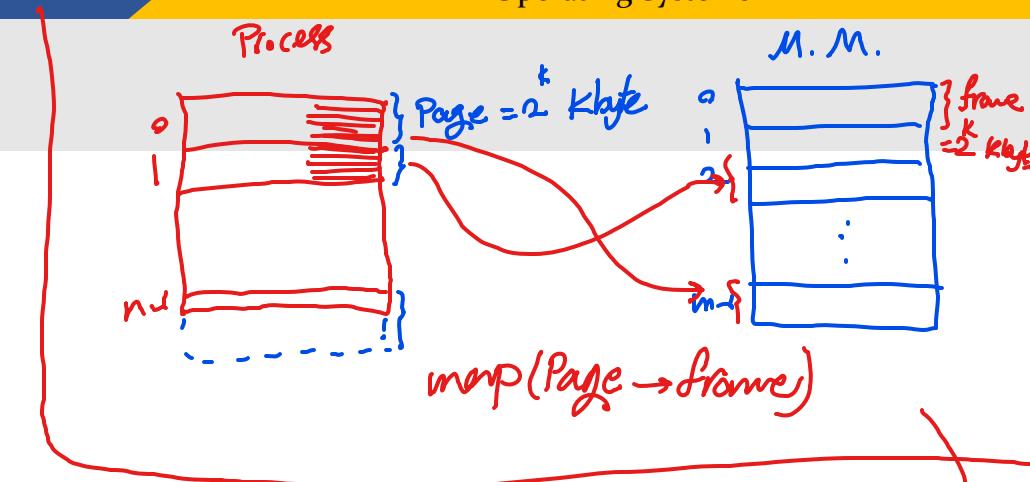
➤ Process is allocated physical memory whenever physical memory is available

- Avoid external fragmentation
- Avoid problem of varying sized memory chunks

➤ Divide physical memory into fixed-sized blocks called frames.

- Size is power of 2, between 512 bytes to 16 Mbytes

➤ Divide logical memory into blocks of same size called pages.



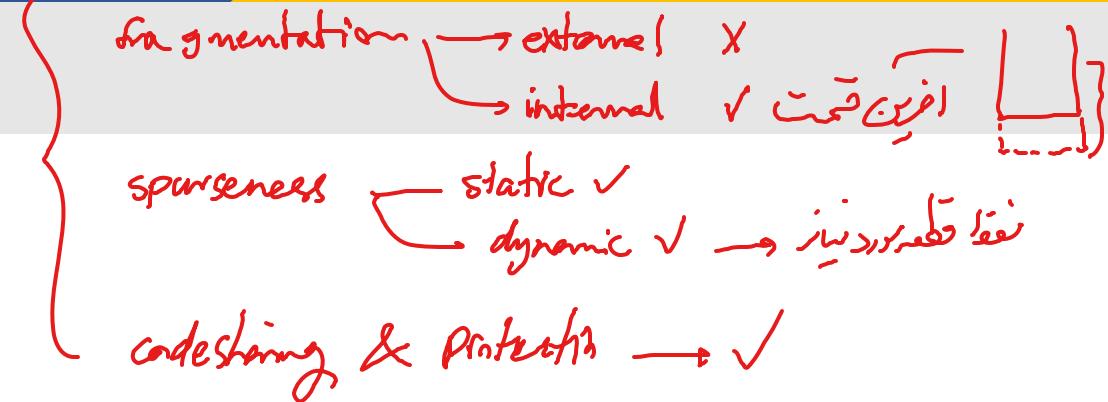
Segment table offset < Page table offset *

جایز است

HARD الحدود ایشان را نیز (حریم) $\text{Page offset} < \text{MM offset}$ *

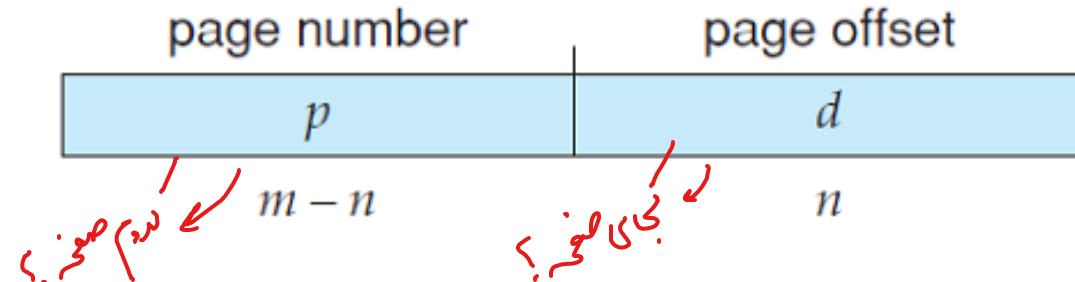
Paging

- Keep tracks of all **free frames**
- To run a program with **N pages**, need to find **N different free frames** and load program.
- Setup a **page table** to translate **logical to physical address**
- Still have **internal fragmentation** *



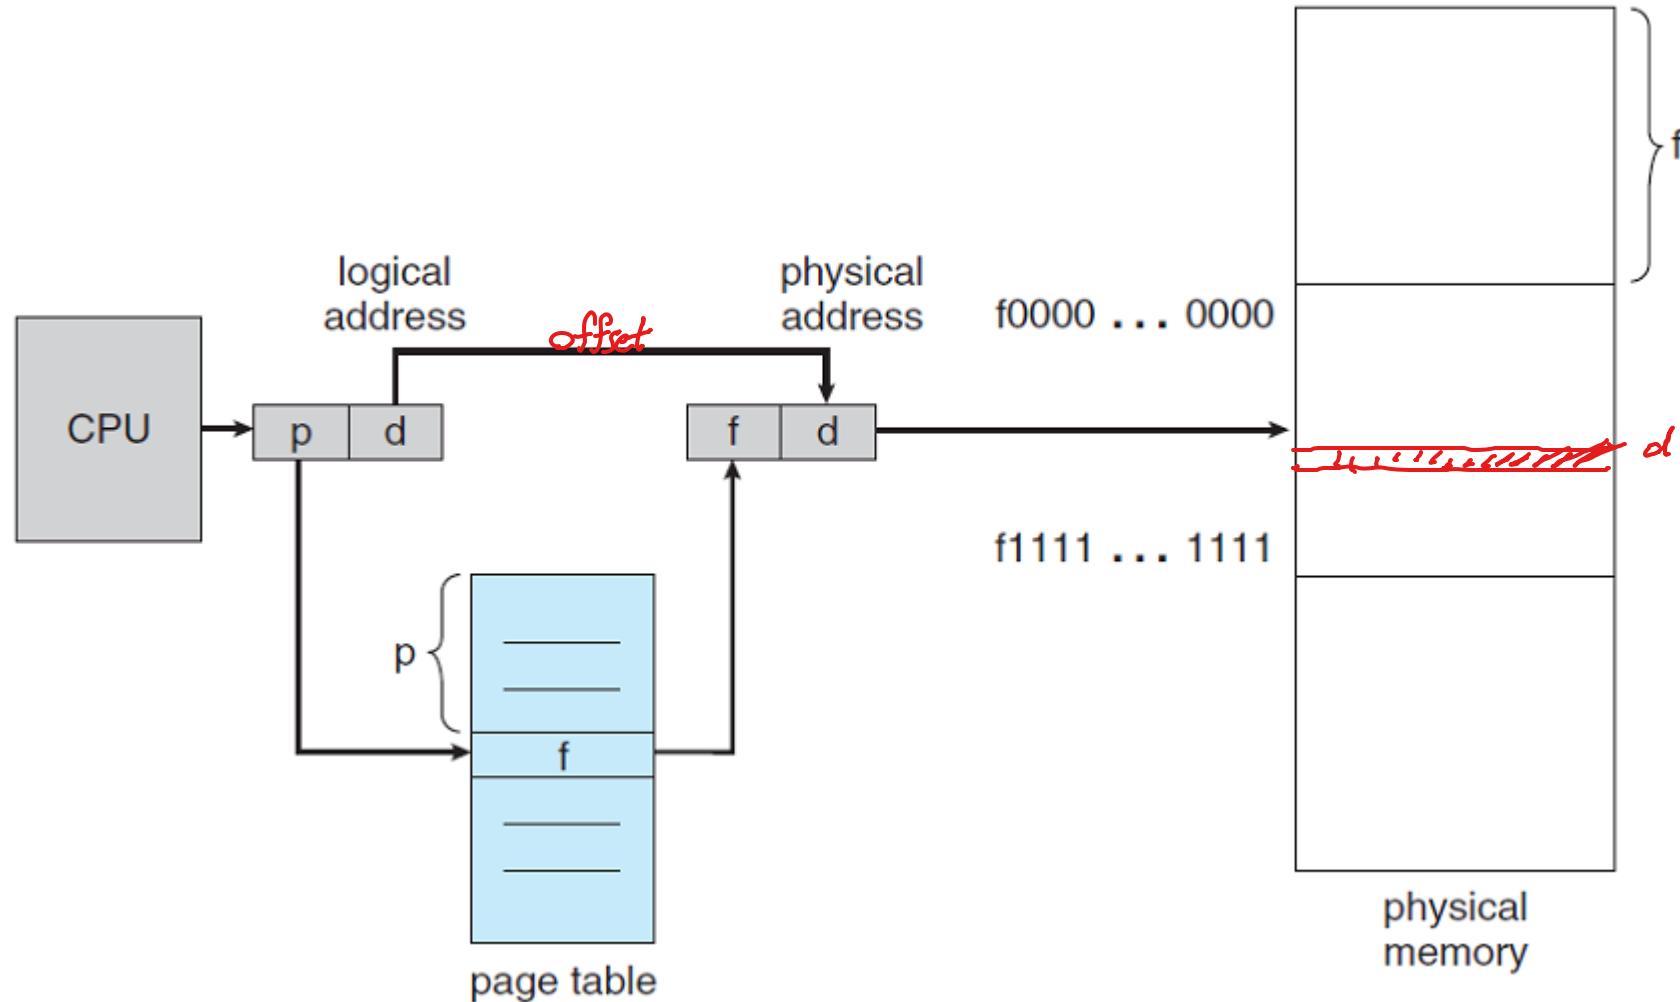
Address translation scheme

- CPU address (Logical address) is divided into two parts:
 - Page number (p): used as an index into a page table that contains base register of each page in physical memory
 - Page offset (d): combined with base address to define the physical memory address that is sent to the memory unit.

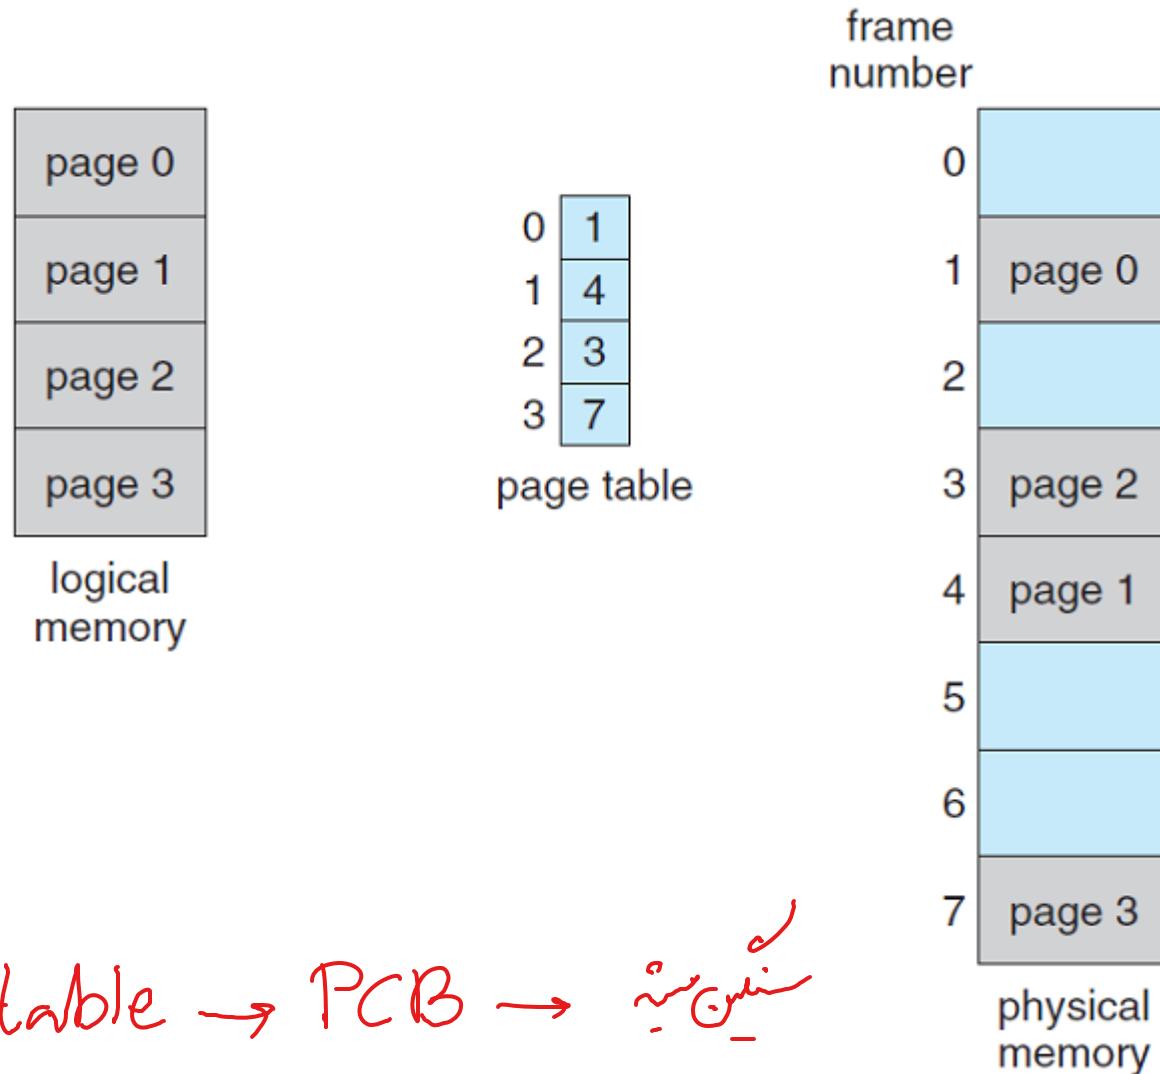


- For given logical address space 2^m and page size 2^n

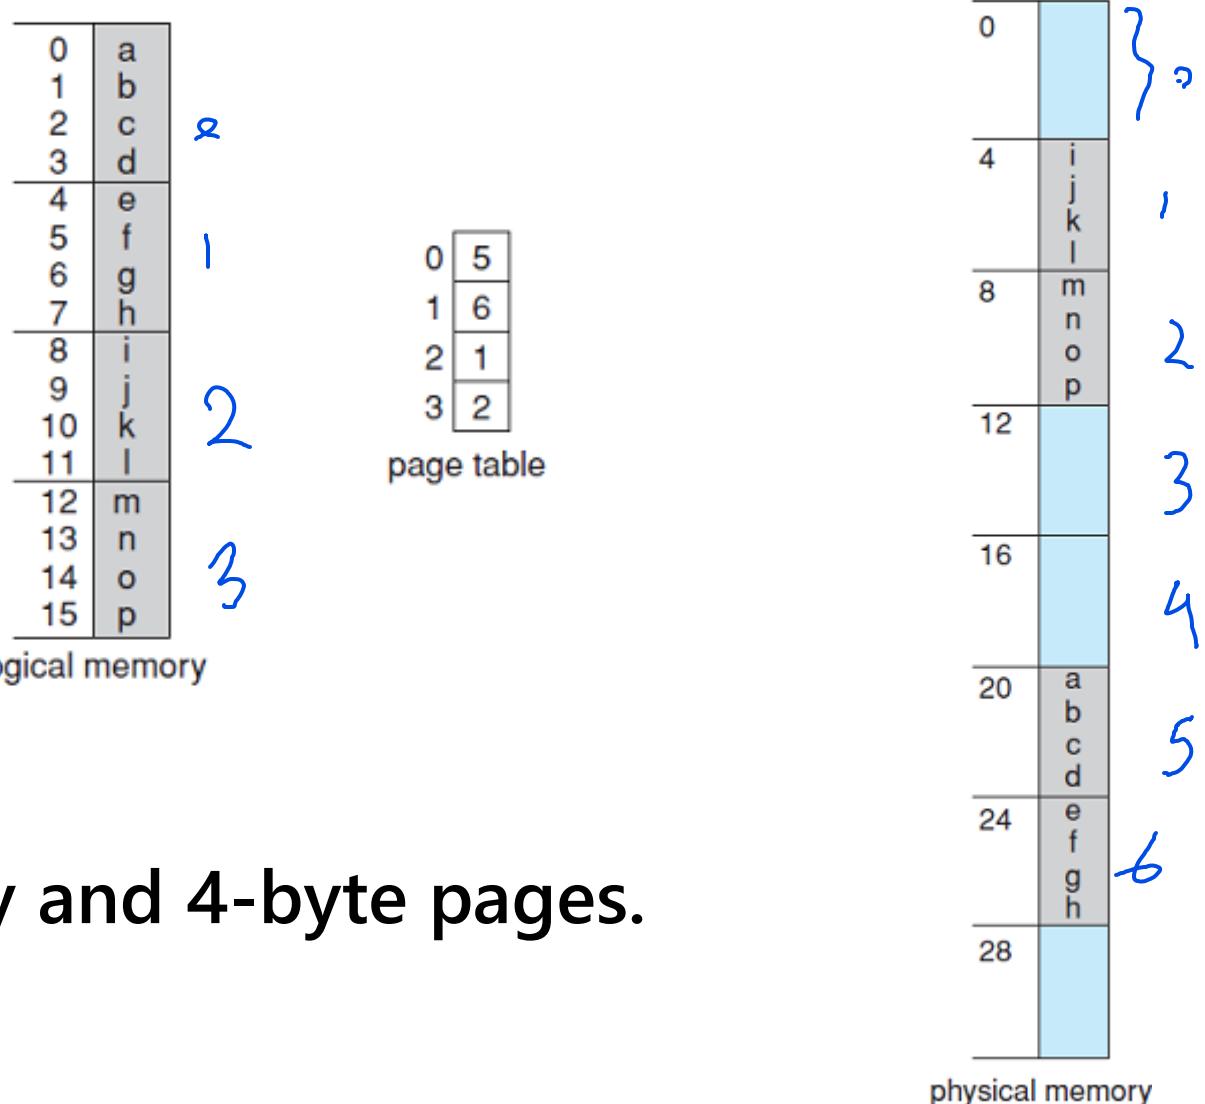
Paging hardware



Paging model of logical and physical memory



Paging example



► **n=2 and m=4, 32-byte memory and 4-byte pages.**

Free frames



Paging example – Internal fragmentation

11 bit offset

➤ Page size = 2048 bytes (2kB)

➤ Process size = 72766 bytes

➤ $72766 / 2048 = 35 \text{ pages} + 1086 \text{ bytes}$

➤ Internal fragmentation: $2048 - 1086 = 962 \text{ bytes}$

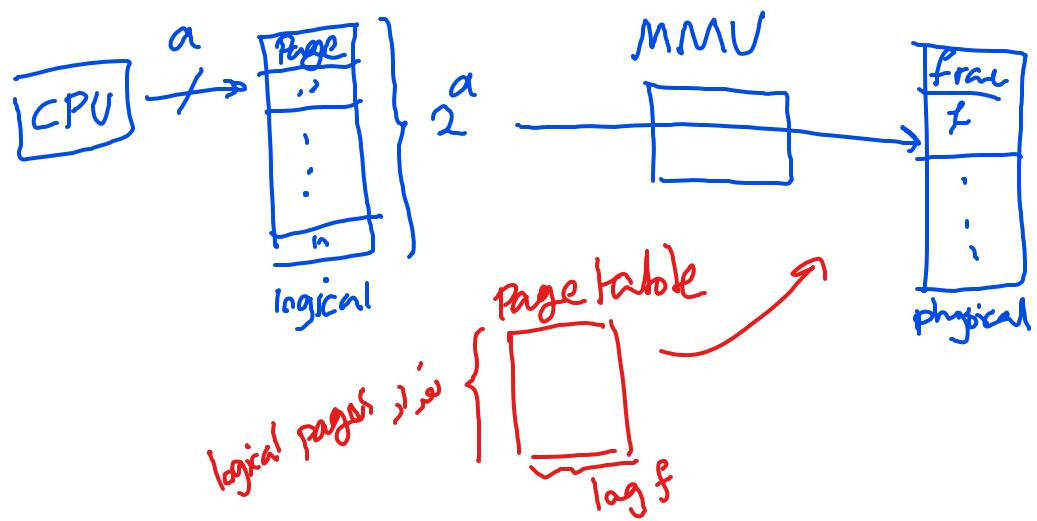
➤ Worst case fragmentation: 1 frame – 1 byte = 2047 bytes

➤ On average fragmentation: $\frac{1}{2}$ frame size

Small page size vs big page size



- On average fragmentation: $\frac{1}{2}$ page size, small page size are good.
- Small page size, more overhead is in the page-table, this overhead is reduced when page size increases.
- Disk I/O is more efficient when the amount of data being transferred is larger (e.g. big pages)
- Page typically are between 4 kB and 8 kB in size.

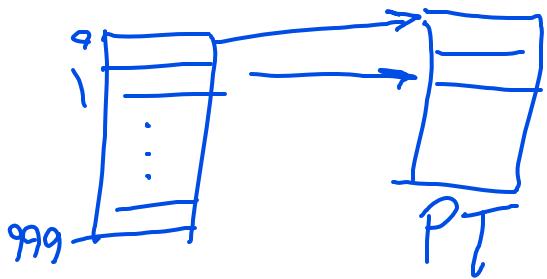


$P \downarrow \rightarrow \text{Page table} \uparrow \Rightarrow \frac{\# \text{ Pages} \times (\log_2 f) \times \frac{1}{8}}{\text{Page size}}$

→ PT *میزهای پیش فرضی*

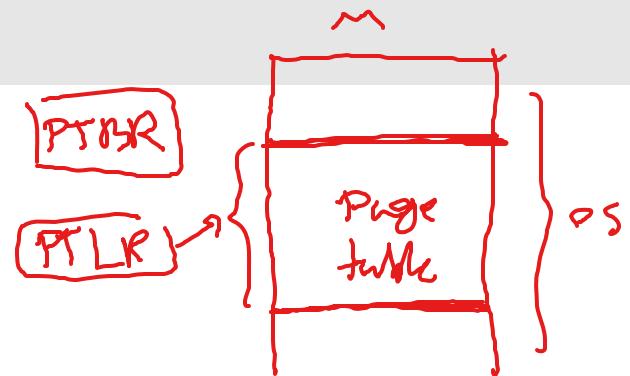
Page table implementation

Size: Page table = 1000



Page table

➤ Page table is kept in memory



➤ Page-table base register (PTBR) points to the page table

➤ Page-table length register (PTLR) indicates size of the page table

➤ In this scheme every data/instruction access requires two memory accesses:

- One for the page table, another for data/instruction



Translation look-aside buffer (TLB)

- The **two memory access problem** can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffer (TLB)**.

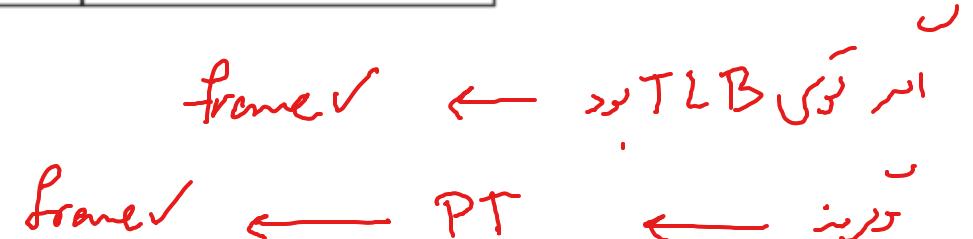
- **Associative memory:** parallel search

Address Locality

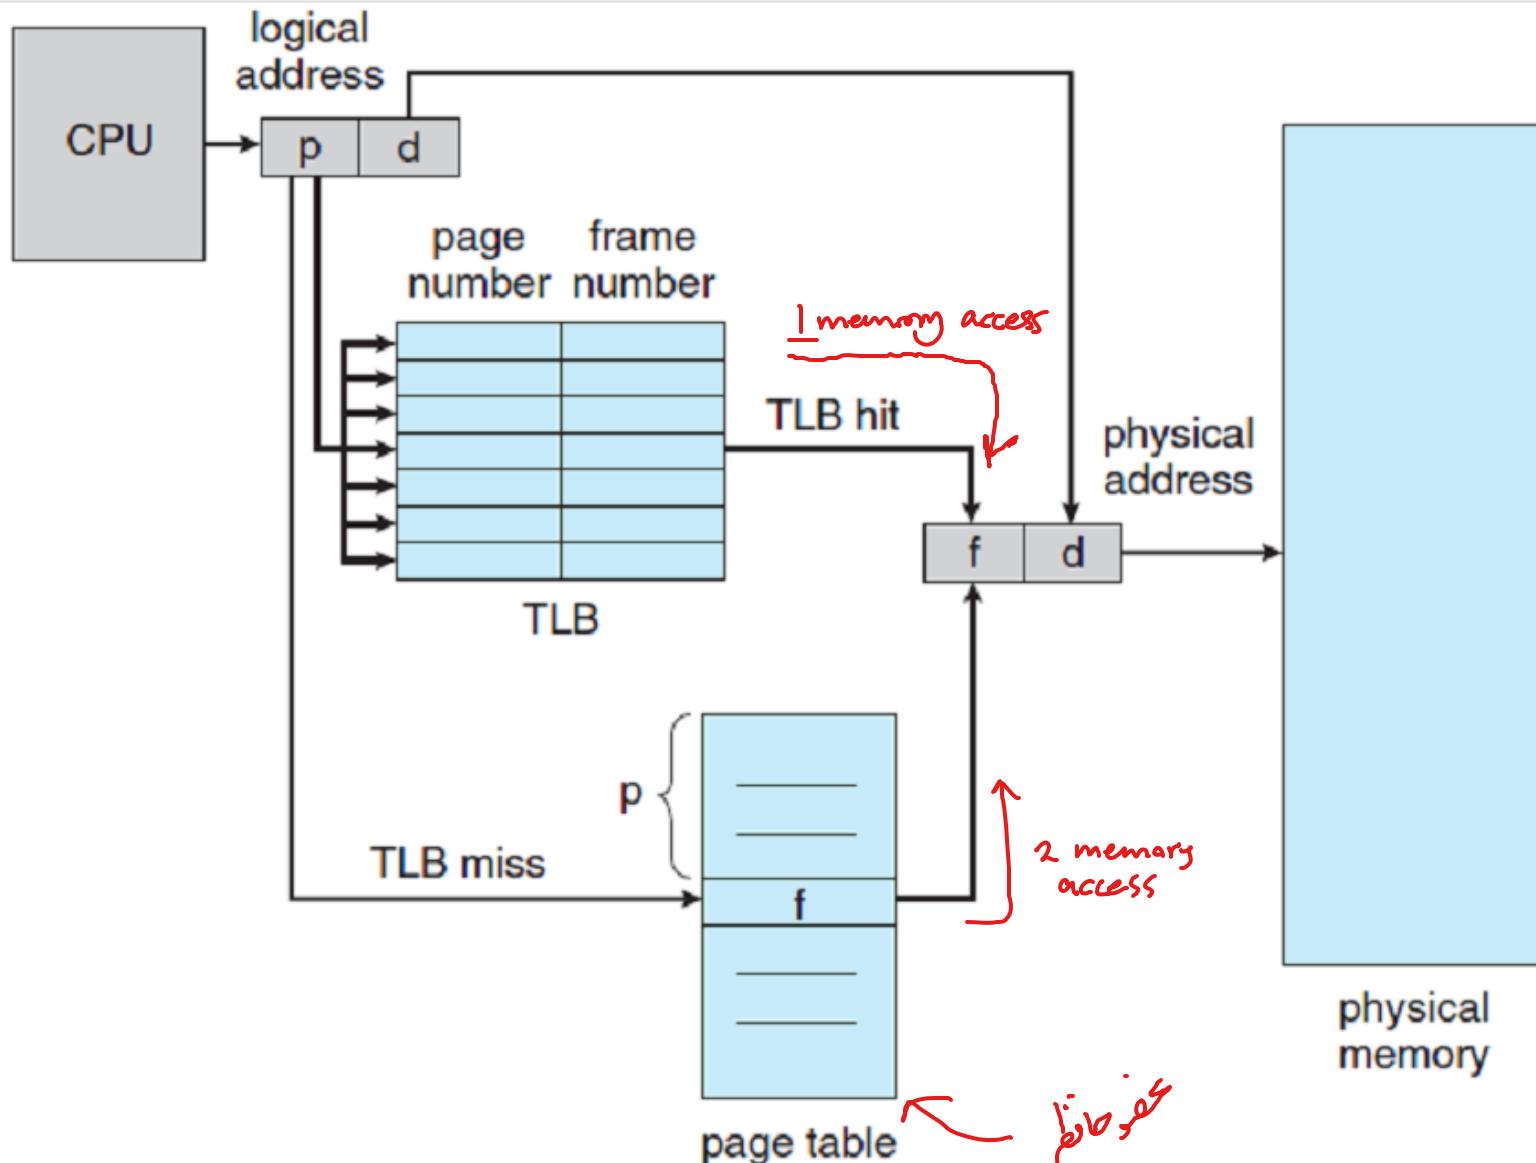
| Page # | Frame # |
|--------|---------|
| | |
| | |
| | |
| | |

Caching ↗ ↘

- **Address translation (p, d):**
 - If p is in associative memory, get frame\# out
 - Otherwise, get frame\# from page table in memory



Paging hardware with TLB



Effective access time

➤ **Hit ratio:** percentage of times that a page number is found in the TLB.

Effective Access Time (EAT)

a = memory access
 b = TLB access
 miss

- α : memory access latency
- h : hit ratio
- $EAT = h \times \alpha + (1 - h) \times 2\alpha$

$$\hookrightarrow h(a+b) + (1-h)(b+2a)$$

$$h = 80\%, \alpha = 100\text{ns} \Rightarrow EAT = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$$

$$h = 99\%, \alpha = 100\text{ns} \Rightarrow EAT = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$$

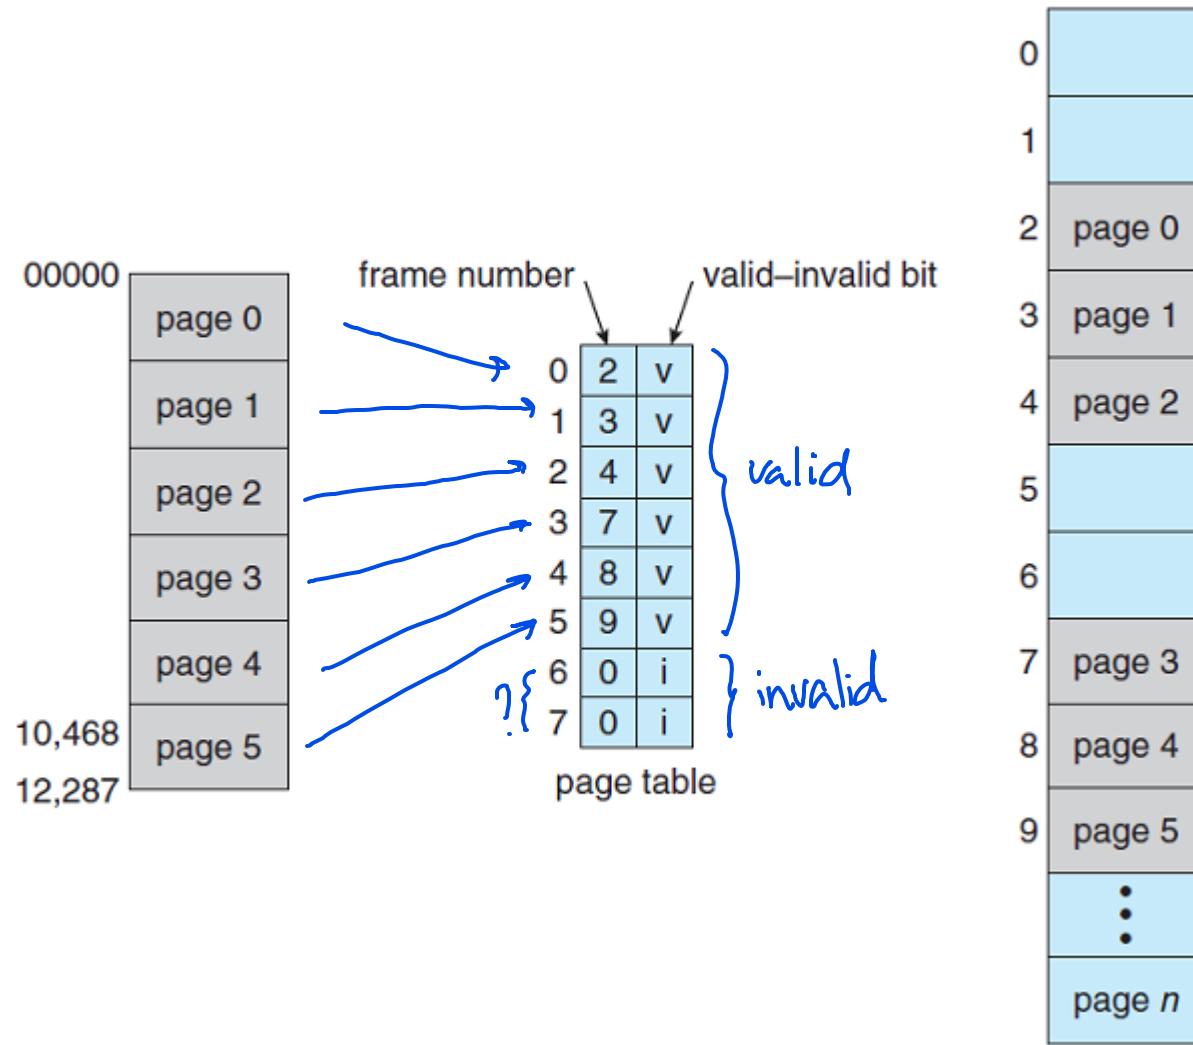
More about TLB

- Some TLBs store **address-space identifier (ASID)** in each TLB entry
 - Uniquely identifies each process to provide address-space protection for that process
 - Otherwise, need to **flush** at every context switch
- TLB is typically **small** (64 to 1024 entries)
- On a **TLB miss**, value is **loaded** into the TLB for faster access next time.
 - Replacement policies must be considered.

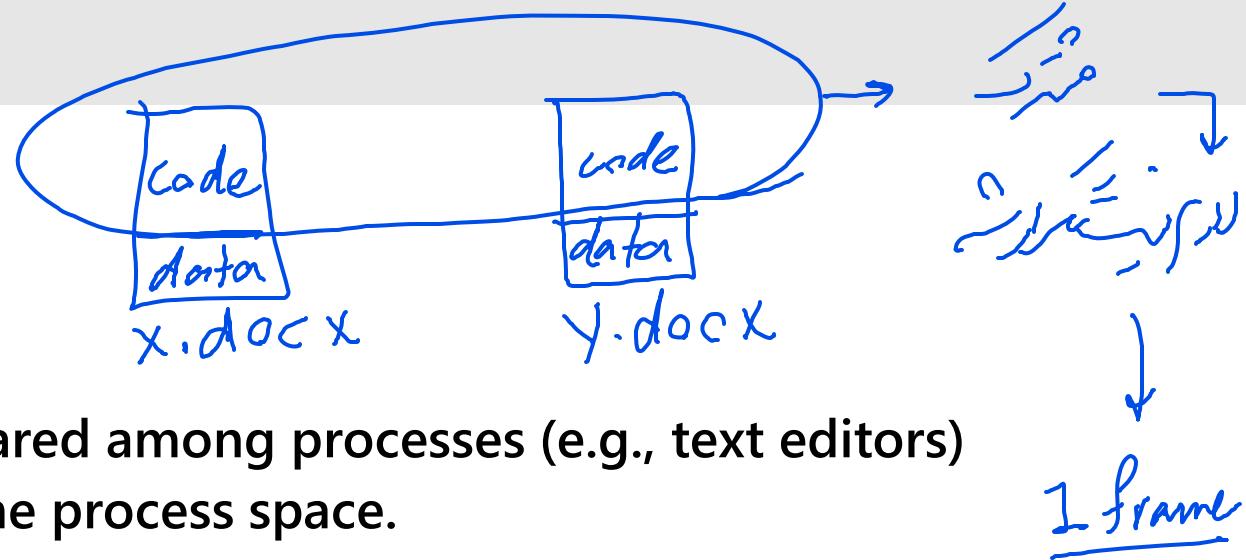
Memory protection

- Memory protection is implemented by **protection bit** with each frame to indicate if read-only or read-write access is allowed.
- **Valid-invalid bit** attached to **each entry** in page table:
 - **Valid** indicates that the page is in the process logical address space (legal page)
 - **Invalid** indicates that the page is not in the process logical address space (illegal page)
 - Or use **page-table length register (PTLR)**
- Any violation result in **trap to the kernel**

Valid/invalid bit in a page table



Shared pages



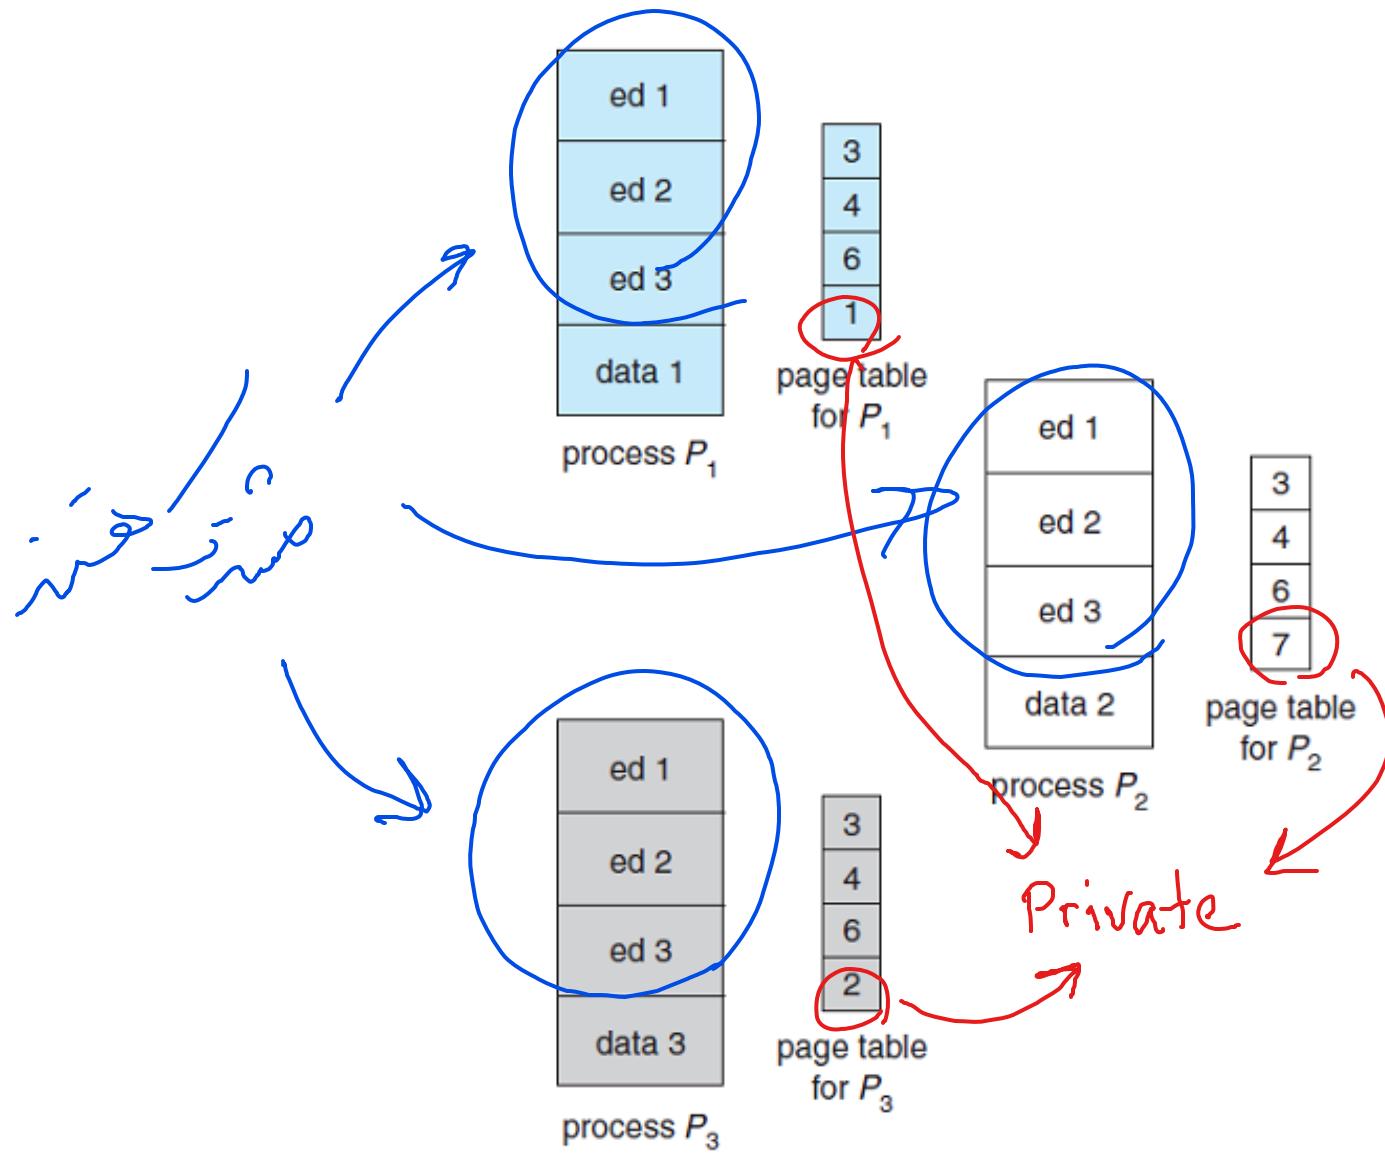
➤ Shared code

- One copy of **read-only (reentrant) code** shared among processes (e.g., text editors)
- Similar to multiple threads sharing the same process space.

➤ Private code and data

- Each process keeps a **separate copy of the code and data**
- The page for the **private code** and data can appear **anywhere** in the logical address space

Shared pages - Example



| | |
|----|--------|
| | 0 |
| 1 | data 1 |
| 2 | data 3 |
| 3 | ed 1 |
| 4 | ed 2 |
| 5 | |
| 6 | ed 3 |
| 7 | |
| 7 | data 2 |
| 8 | |
| 9 | |
| 10 | |
| 11 | |

اجزای اعلیٰ مراجعت
که رفع نموده ام
برای اینجا اینجا

Problem of big page tables

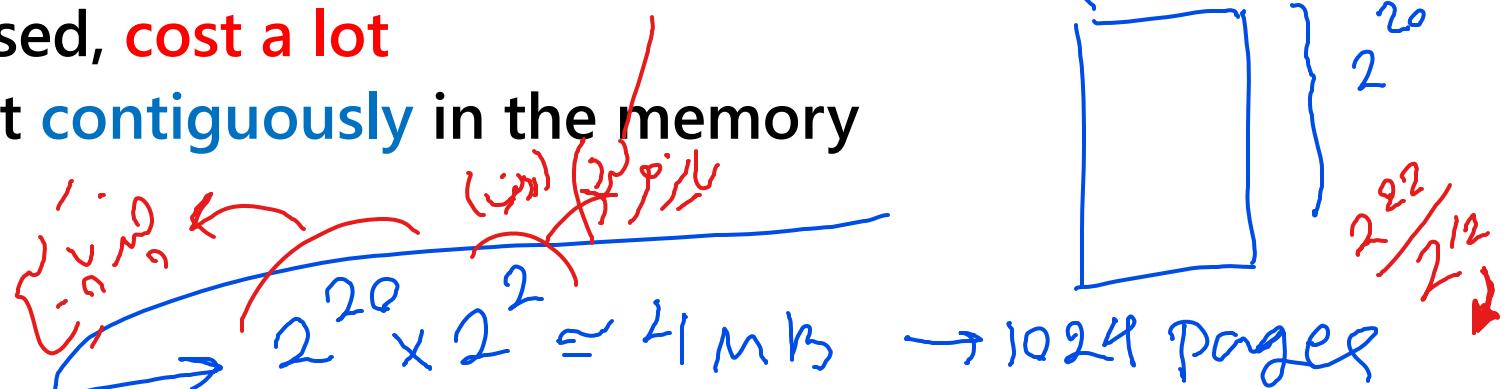
← { Page size ↓
 ↗ ↘ ↑

➤ Memory structure for paging can get **huge** using straight-forward methods.

$$\xrightarrow{32} 2^{32}$$

➤ Consider a 32-bit logical address space on a modern computers:

- Page size of **4 kB** = 2^{12}
- Page table would have 1 million entries ($2^{32} / 2^{12}$) = 2^{20} **# Page (PT [i])**
- If each entry is **4 B**: **4 MB** of physical address space memory for page table alone.
- The amount of memory used, **cost a lot**
- **Don't** want to allocate that **contiguously** in the memory

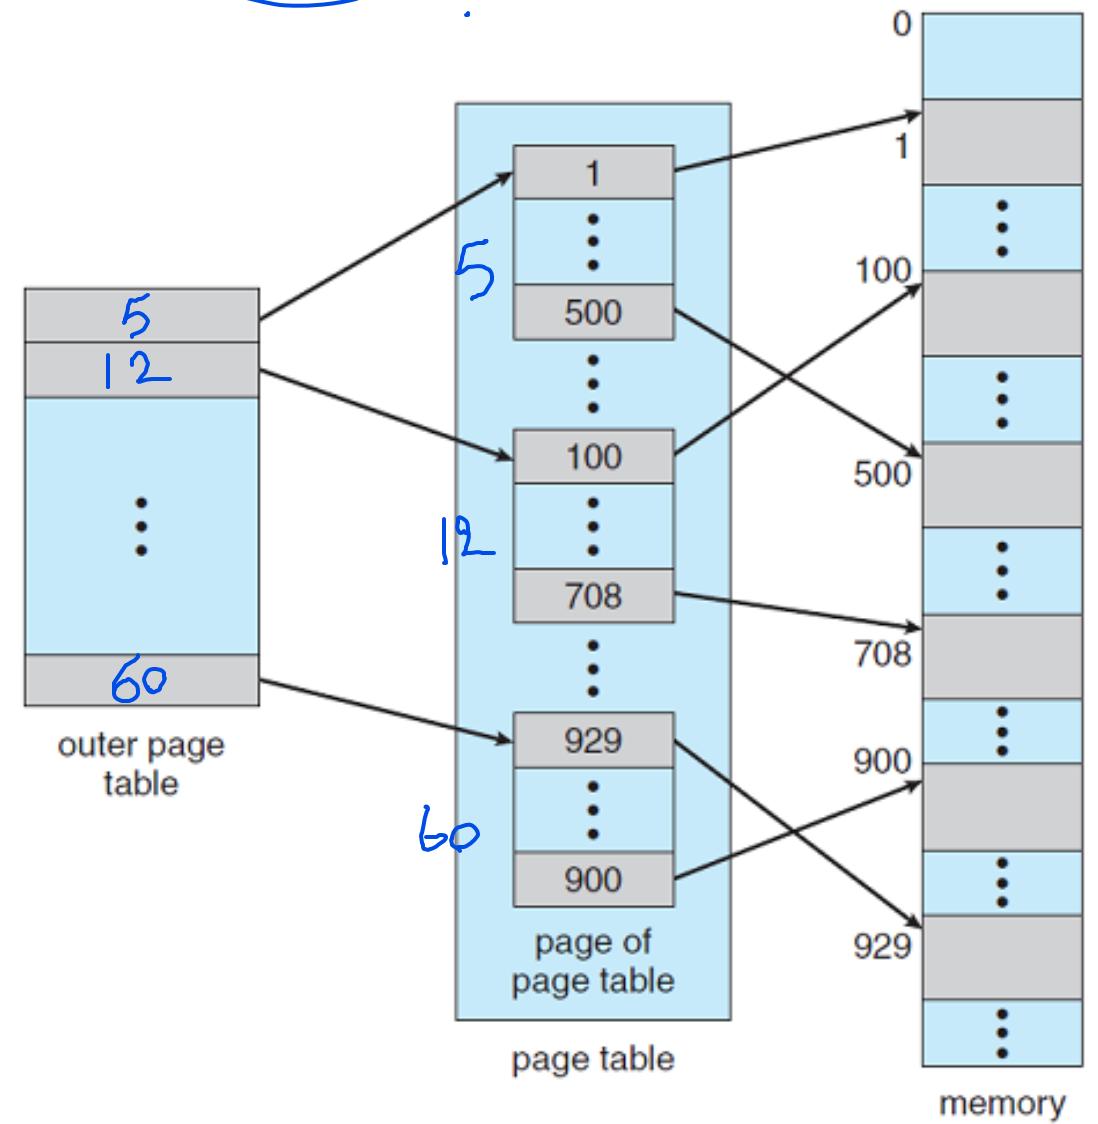


Solutions to maintain **huge** page tables

- A) **Hierarchical paging**
- B) **Hashed page tables**
- C) **Inverted page tables**

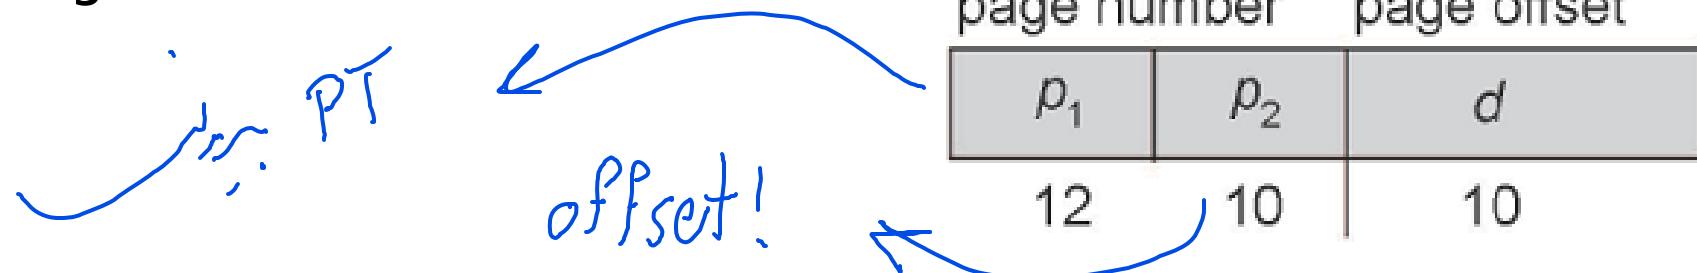
A) Hierarchical paging

- Use of two-level page table!
- We then page the page table!



Two-level paging - example

- A logical address, on 32-bit machine with 1 kB page size, is divided:
 - A page offset consisting 10 bits
 - A page number consisting of 22 bits
- Since the page table is paged, the page number is divided into:
 - A 12-bit page number
 - A 10 bit page offset
- Thus a logical address is:

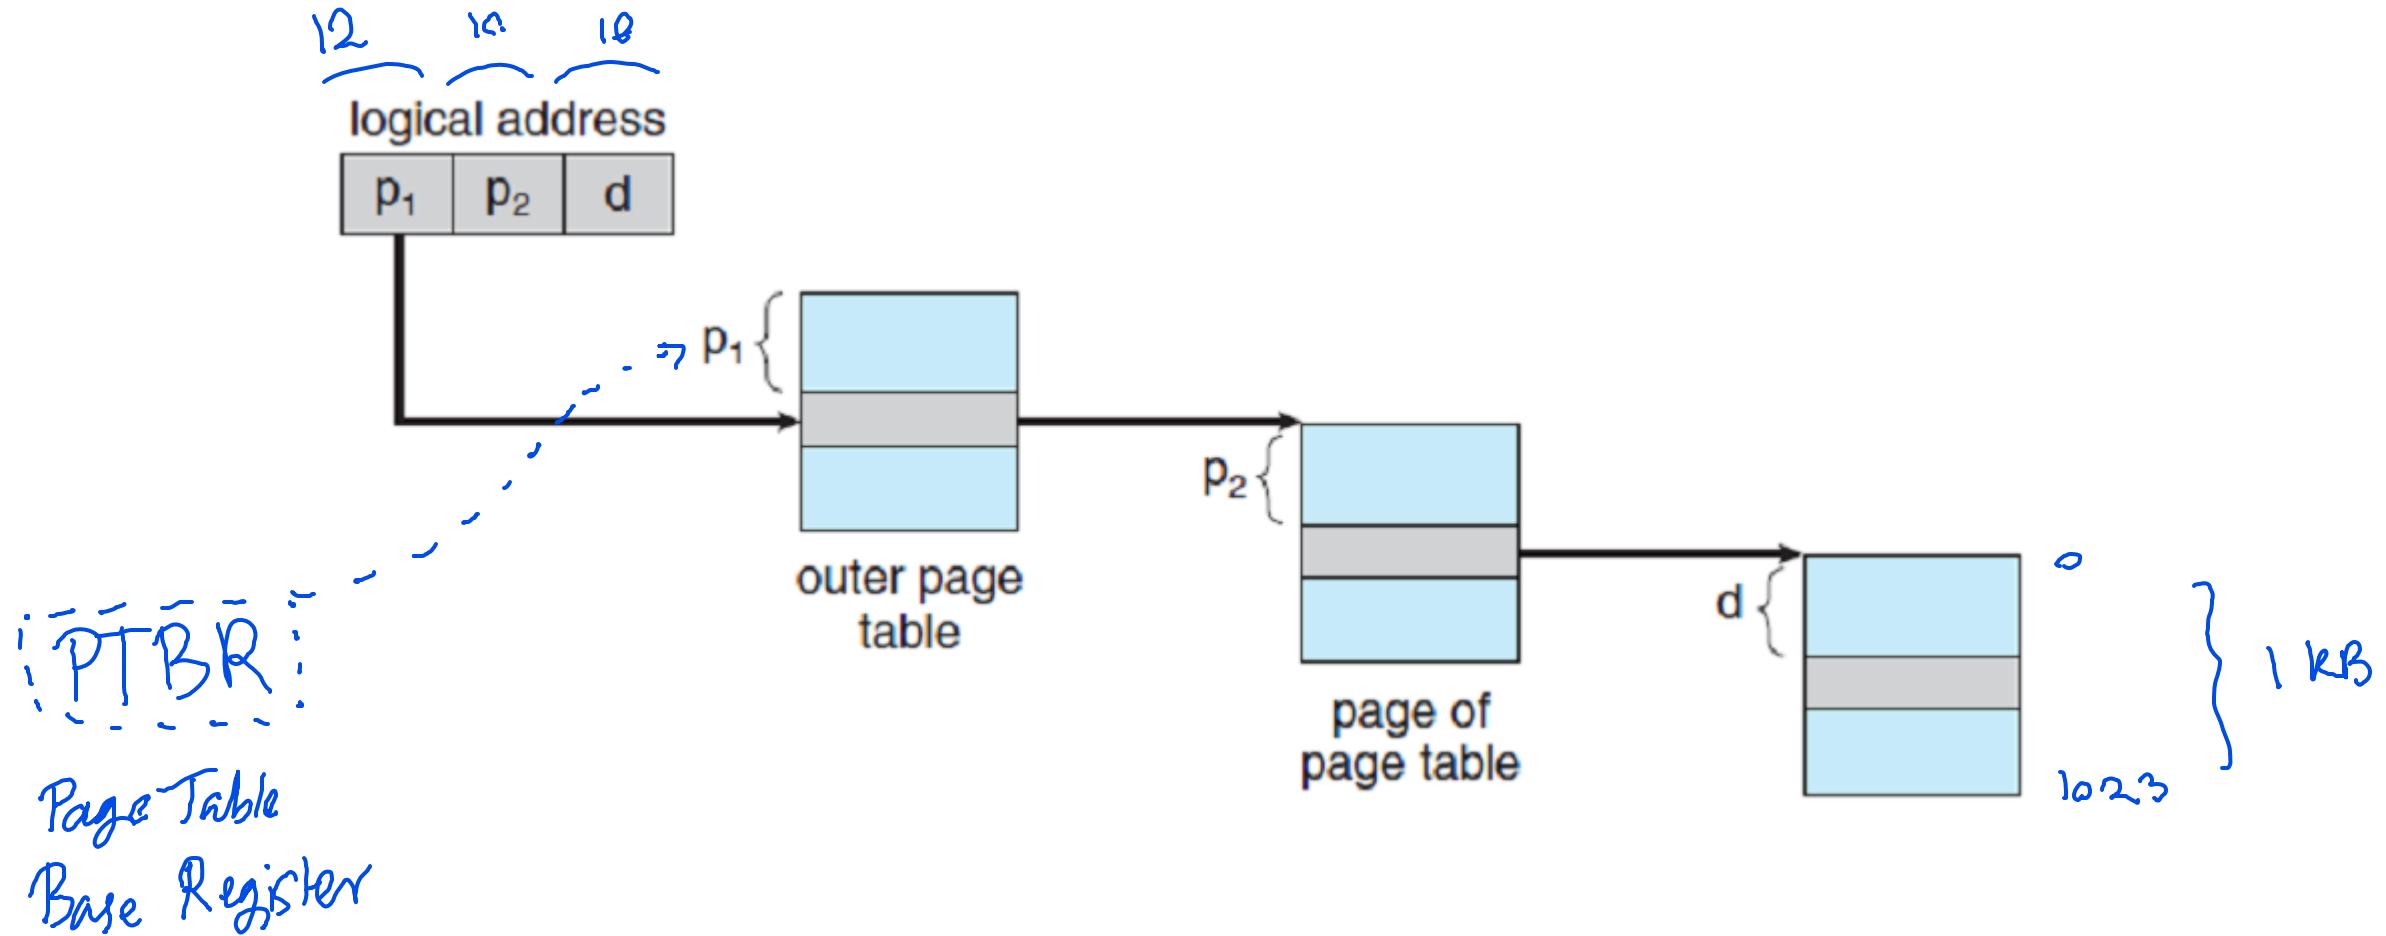


$$\frac{2^{32}}{2^{10}} = 2^{22}$$

where p_1 is an index into outer page table, and p_2 is the displacement within the page of inner page table

- Known as **forward-mapped page table**

Address-translation scheme



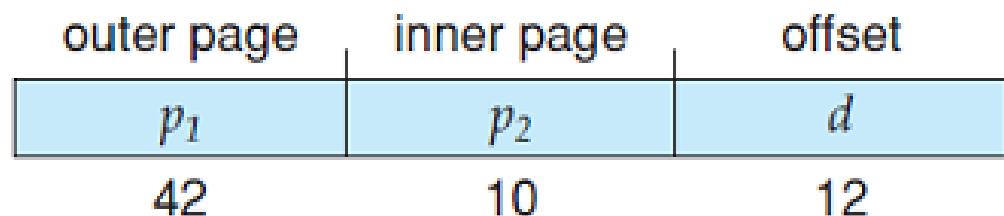
64-bit logical address space

➤ Even two-level paging scheme is not sufficient



➤ If page size of 4 kB (2^{12})

- Page table entries = 2^{52}
- Assuming each page table entry size = 4 B
- If two-level scheme, inner page table could be 2^{10} , 4 kB entries
- Outer page table has 2^{42} entries or 2^{44} B
- Address would look like:



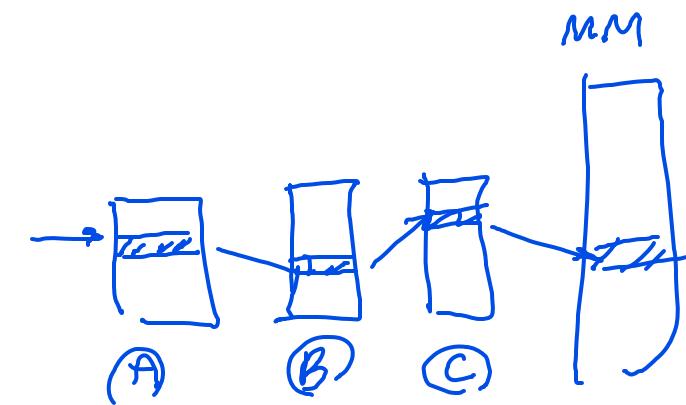
Three-level paging scheme

- One solution is to add a **2nd outer page table**
- But in the following example, the 2nd outer page table is still **2^{34} bytes** in size!
- And possible **4 memory access** to get to one physical memory location:

| outer page | inner page | offset |
|------------|------------|--------|
| p_1 | p_2 | d |
| 42 | 10 | 12 |

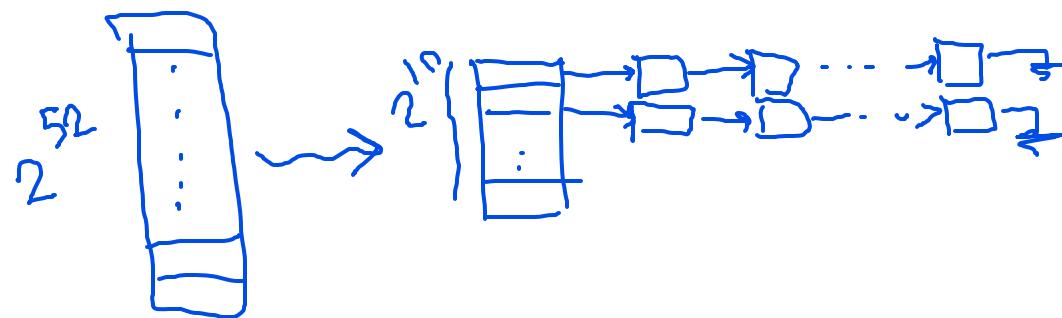
| 2nd outer page | outer page | inner page | offset |
|----------------|------------|------------|--------|
| p_1 | p_2 | p_3 | d |
| 32 | 10 | 10 | 12 |

(A) (B) (C)



B) Hashed page tables

- Common in address space > 32 bits
- The logical page number is hashed into a page table
- This page table contains a chain of elements hashing to the same location



Hash page table scheme

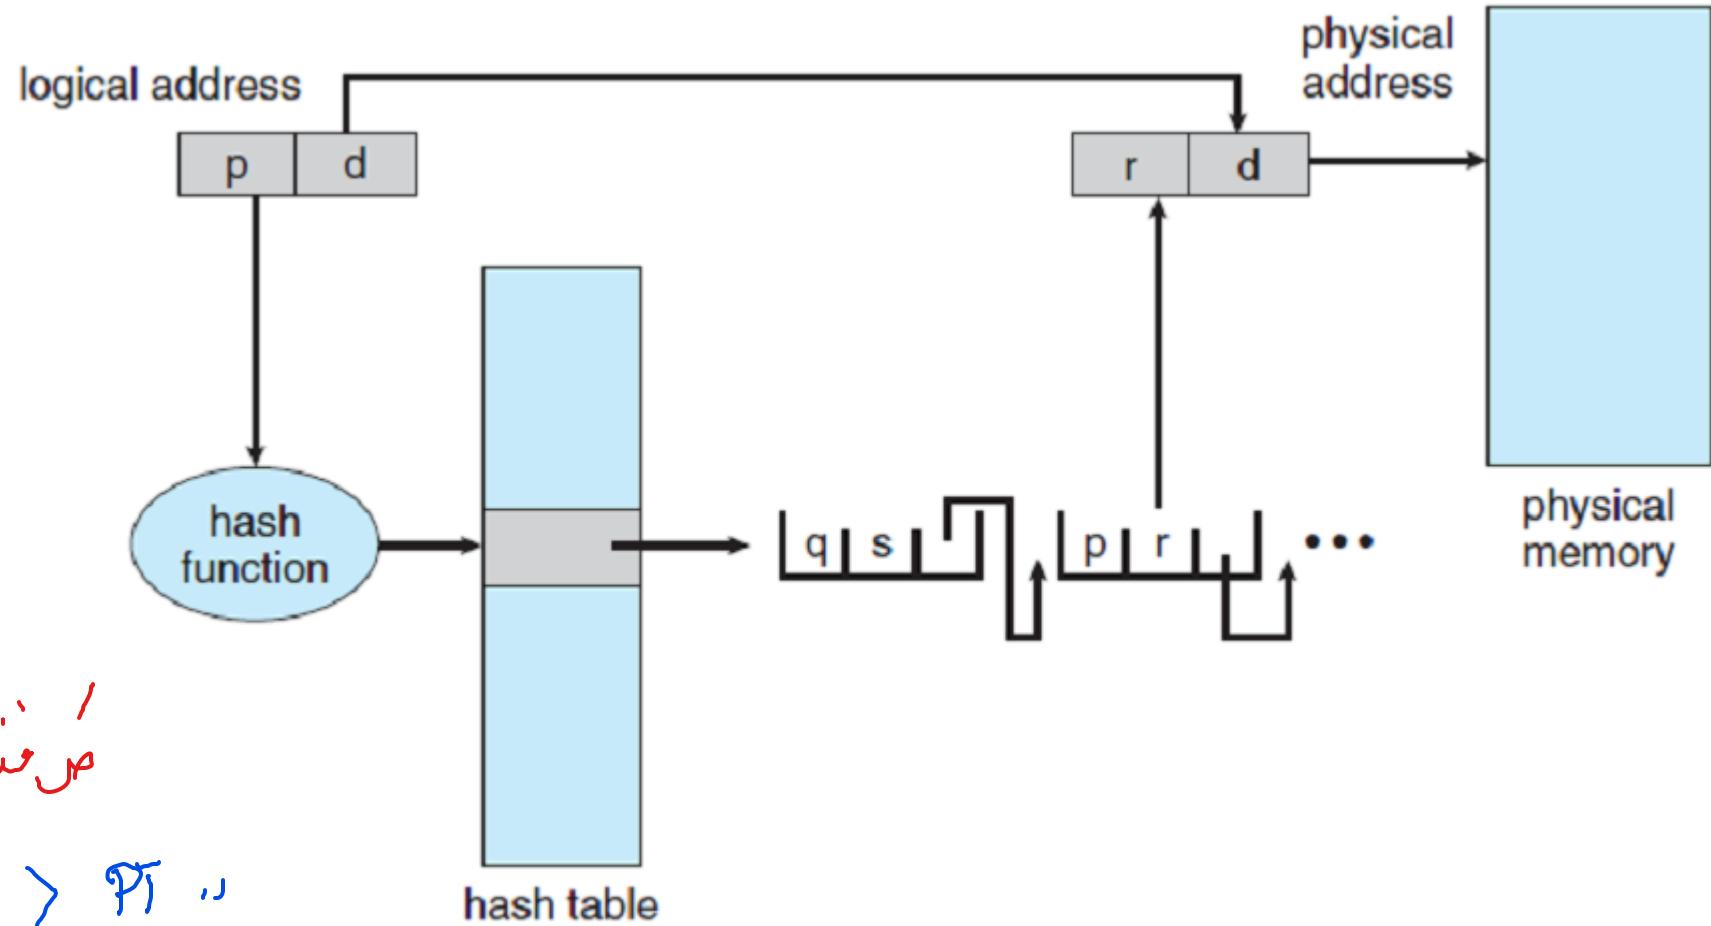
➤ Each element contains

- Logical page number
- Physical frame number
- Pointer to the next node

➤ Search is done **serially** in the **linked list**

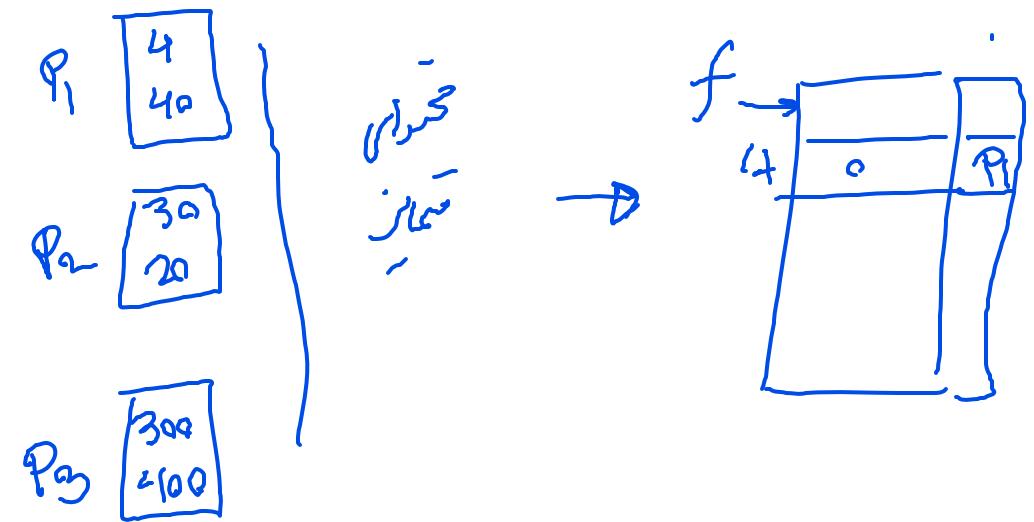
Handwritten notes:

- HashT → PT
- HT > PT
- HT \rightarrow PT
- مسیرهایی از صفحه هایی که در حافظه ممکن است قرار گیرند
- مسیرهایی از صفحه هایی که در حافظه ممکن است قرار گیرند
- ptr

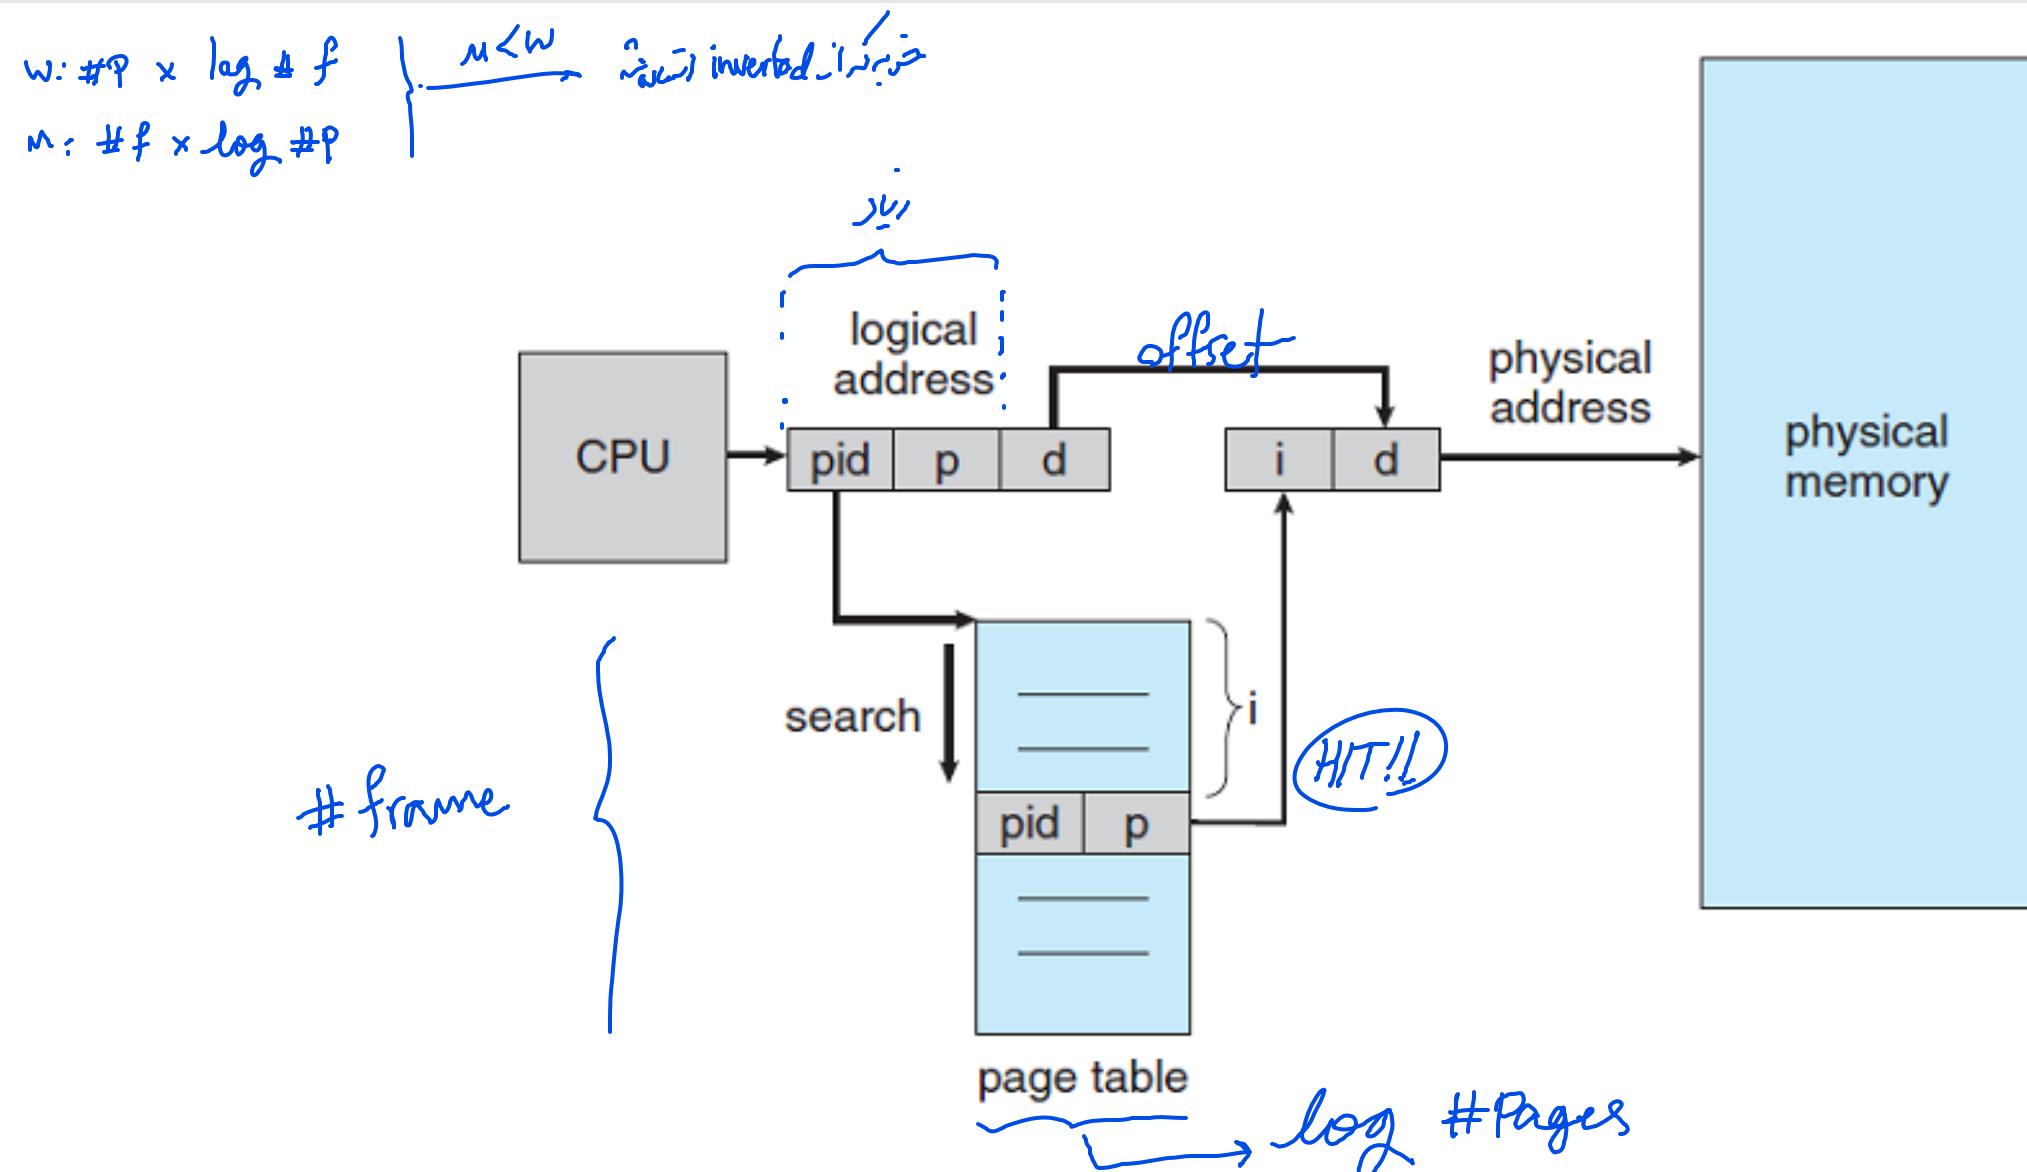


C) Inverted page table

- Rather than keeping all **possible logical page numbers**, track all **physical pages (frame)** numbers
- One entry for each physical page of memory
- Entry consists of
 - **Virtual address** of the page stored in that frame
 - + Process info



Inverted page table scheme



Problem and solutions

➤ Good:

- Decrease memory needed

➤ Bad:

- Increase time needed to search the table

➤ Use **hash table** to limit the search to one, or at most a few, page-table entries.

➤ How to implement **shared memory**?

- One mapping of a virtual address to the **shared physical address**

Questions?

